

**SUPSI**

# Computer Graphics

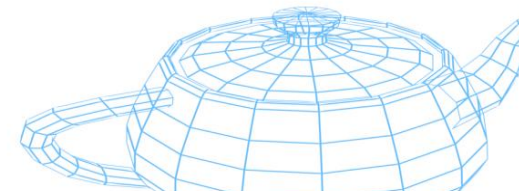
OpenGL (1): FreeGLUT, contexts, buffers and first steps

Achille Peternier, lecturer



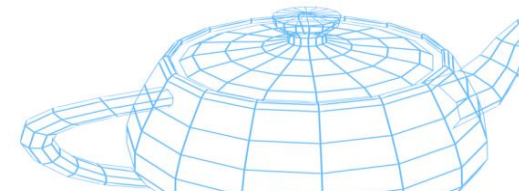
## FreeGLUT

- **Free** alternative to the Open**GL** **U**tility **T**oolkit.
- Evolution/clone of GLUT, originally created by Mark Kilgard (Nvidia) and abandoned in 1999:
  - Less restrictive license.
- FreeGLUT started in 1999 by Pawel Olzsta:
  - Open-source.
  - No code inherited from GLUT.



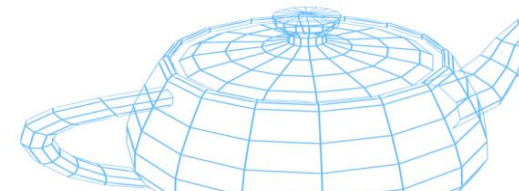
## FreeGLUT

- Event-based (callbacks).
- Supports user-input through mouse and keyboard.
- Provides a menu/sub-menu system.
- Supports printing text to the OpenGL window.
- Provides a series of built-in, dynamically generated 3D objects.
- Available on Windows, Linux, MacOS, etc.



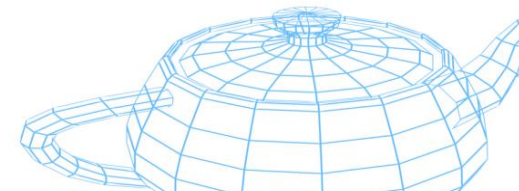
## FreeGLUT

- Windows:
  - Download and compile the project:
    - <http://freeglut.sourceforge.net/>
  - Use the .lib + .dll or .lib-only (static) version.
  - Define **FREEGLUT\_STATIC** for static linking:
    - The static lib is included (and used) in the tutorials and series' solutions.
- Ubuntu:
  - Execute: `sudo apt install freeglut3-dev`
  - Link to *libglut.so* (dynamic) or *libglut.a* (static).




# FreeGLUT

- Typical usage:
  - Initialize the library.
  - Create one or more windows for graphic output.
  - Create menus (optional).
  - Register callback functions.
  - Enter the main loop.



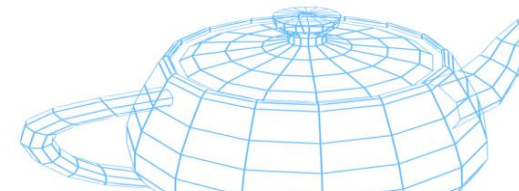
## FreeGLUT

- `void glutInit(int *argc, char *argv[]);`
  - Accepts -display, -geometry
- `void glutInitDisplayMode(flags);`
  - Accepts:
    - `GLUT_SINGLE, GLUT_DOUBLE` single/double buffering
    - `GLUT_RGB, GLUT_RGBA` color mode
    - `GLUT_DEPTH` enables the z buffer
    - ...
  - E.g.: `glutInitDisplayMode(GLUT_DOUBLE | GLUT_RGBA);`




## FreeGLUT

- `void glutInitWindowSize(int width, int height);`
- `void glutInitWindowPosition(int x, int y);`
- `int glutCreateWindow(char *name);`
  - name = window title.
  - returns the window ID.
- `int glutCreateSubWindow(int parent_window, int x, int y,  
int width, int height);`



## FreeGLUT

- Callback registration:
  - `glutDisplayFunc(func. ptr.);`
    - Invoked each time the output scene must be re-rendered.
    - Call `glutPostWindowRedisplay(winId)` to force a refresh.
  - `glutReshapeFunc(func. ptr.);`
    - Triggered each time the window size is changed.
  - `glutMouseFunc(func. ptr.);`
    - Triggered each time the mouse is used.
  - `glutKeyboard(func. ptr.);`
    - Triggered each time a keyboard key is pressed.
  - `glutSpecial(func. ptr.);`
    - Same as before, but for special keyboard keys such as the arrows.








# Tutorial

Basic FreeGLUT

## RGB




- **Red Green Blue :**


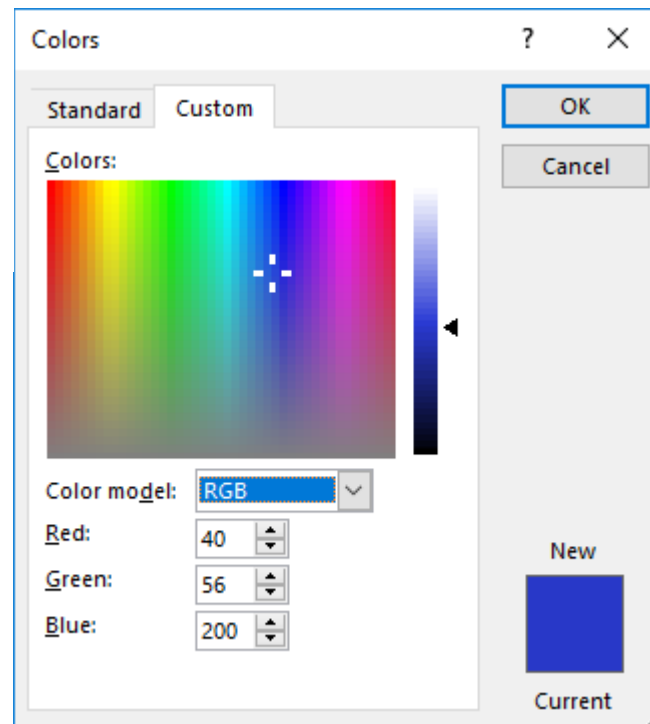
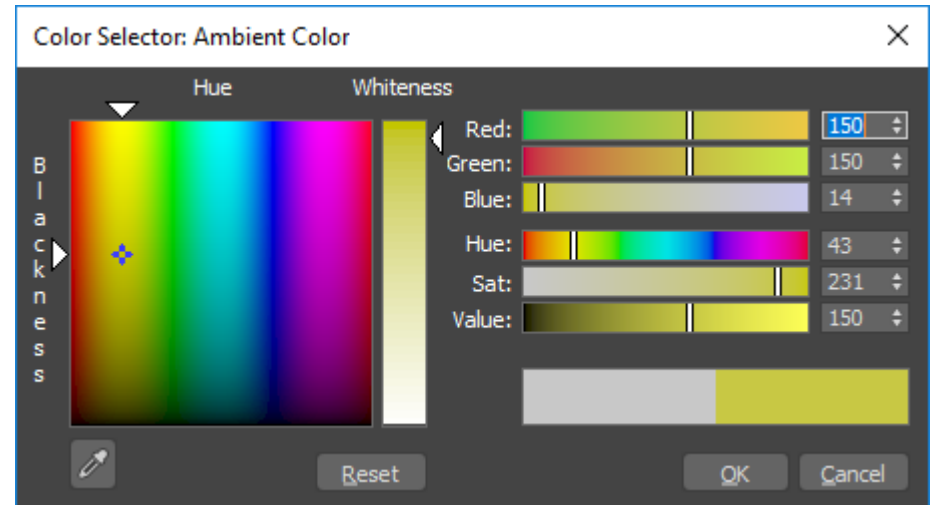
- Color model used to express colors based on the intensity of the three base colors.
- Works for light-emitting sources (and not for painting).
- Additive method:

$$\text{color} = R_{\text{intensity}} + G_{\text{intensity}} + B_{\text{intensity}}$$

- Several ways to encode values:


- Bytes [0-255], e.g.: **[255 0 0]**, **[0 255 0]**, [128 128 128]
- Float [0.0-1.0], e.g.: **[1.0 0.0 0.0]**, **[0.0 1.0 0.0]**, [0.5 0.5 0.5]
- Hexadecimal [00-FF], e.g.: **#FF0000**, **#00FF00**, #7F7F7F





## RGBA

- Another channel (alpha) is added for storing additional information (**usually** transparency):
  - As already seen for the intensities of the other channels, the alpha channel intensity is defined as *0 = transparent*, and *max value = completely solid*.
- RGBA using float =  $4 * \text{sizeof(float)} = 16 \text{ bytes} = 128 \text{ bit}$ :
  - Good for memory alignment.
  - Perfect for SIMD 128 bit registers.
- Specifying transparent alpha values **will not** automatically activate transparency in OpenGL!
  - Transparency is a much more complex topic that we will see later.




## RGBA

- No need to create a new class: RGBA = XYZW.
  - Reuse `glm::vec3`, e.g.:

```
// Define red [1.0 0.0 0.0]:  
glm::vec3 color;  
    color.r = 1.0f;  
    color.g = 0.0f;  
    color.b = 0.0f;
```

- ...or `glm::vec4` for RGBA colors, e.g.:

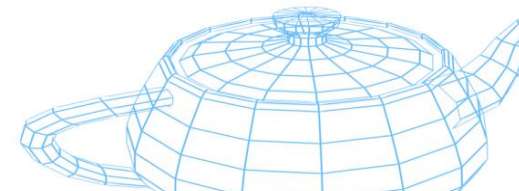
```
// Define red with alpha channel [1.0 0.0 0.0 1.0]:  
glm::vec4 color;  
    color.r = 1.0f;  
    color.g = 0.0f;  
    color.b = 0.0f;  
    color.a = 1.0f;
```






## Main buffers

- A rendering-context is initialized by specifying the characteristics of these four buffers:
  - **Framebuffer** = main output buffer (where you render the final image).
  - **Back buffer** = the hidden twin of the framebuffer, where the next frame is being rendered [almost mandatory].
  - **Depth buffer** (or **Z buffer**) = stores Z values for each pixel of the framebuffer [optional].
  - **Stencil buffer** = additional buffer for per-pixel logic operations [really optional].
  - **Accumulation buffer(s)** = one or more additional buffers for storing a series of images [extremely optional].



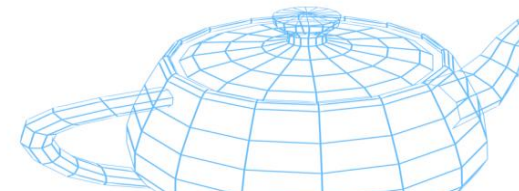

## Framebuffer

- Memory segment containing the image information to display through the graphics device.
- Contains pixel colors:
  - RGB and RGBA on modern displays.
  - Single bits for older, monochromatic monitors.



## Framebuffer

- The memory information stored in the framebuffer is accessed by the device to periodically refresh the pixels rendered on the screen:
  - Typical refresh-rate between 60 and 120 Hz:
    - 120 Hz useful for stereographic rendering (60 Hz per eye).
- To avoid visual artifacts (screen flickering and tearing), video memory refreshing is synchronized with the screen refresh rate.





Tear Point #1 --->

Tear Point #2 --->






## Framebuffer

- Vertical synchronization:
  - Old terminology used on CRT monitors:
    - The graphics card waits for the vertical beam to reach the lower-right corner of the screen.
    - While the beam is reset to its upper-left position, the graphics memory is refreshed.
    - Additional modifications are prevented until the next reset.
  - It is used to synchronize the rendering speed with the monitor refresh frequency.
- Double buffering:
  - Uses two buffers: front (main one, rendered on the screen) and back.
  - Instead of rendering to the screen memory buffer, a secondary (hidden) buffer is used.
  - Once the image is ready on the back-buffer, it is copied to the front-buffer (usually during vertical sync):
    - As optimization, back- and front-buffer pointers are swapped (zero copy, page flipping/ping-pong buffering).
- Triple buffering:
  - Same as double buffering, but two back-buffers are used.
  - With double buffering, when the back-buffer image is ready, the pipeline is stalled waiting for the front-buffer to be available:
    - With triple buffering, there's always a non-locked buffer for rendering.




## Framebuffer

- New technologies for getting rid of the GPU-display synchronization problem:
  - NVidia G-Sync: dedicated HW embedded in monitors to dynamically sync with the GPU signal.
  - AMD FreeSync: a similar mechanism but working without dedicated components and released patent-free.



# Framebuffer

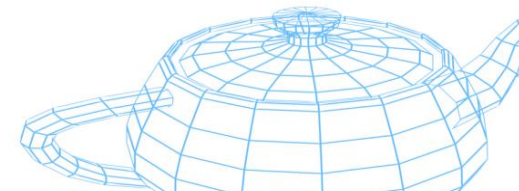
- Different framebuffer color depths:
  - RGB:
    - 1 bit = monochromatic.
    - 4 bit = 16 colors (from a palette).
    - 8 bit = 256 colors (from a palette).
  - RGB/RGBA:
    - 15 bit = high-color ( $2^{15}$  colors):
      - 5+5+5 (RGB).
    - 16 bit = high-color ( $2^{16}$  colors):
      - 5+5+5+1 (RGBA) = 15 bit high-color + alpha channel.
      - 5+6+5 (RGB) = 16 bit high-color:
        - » Human eye is more sensitive to green.
      - 4+4+4+4 (RGBA).
    - 24 bit = true-color ( $2^{24}$  colors, ~16 millions):
      - Human eye can recognize up to 10 million colors.
  - RGBA:
    - 32 bit = true-color (24 bit) + 8 bit alpha channel.
  - > 32 bit:
    - High Dynamic Range (HDR), professional devices, intermediate steps.






## Framebuffer

- RGB or RGBA mode must be specified during the context creation:
  - `glutInitDisplayMode(GLUT_RGB or GLUT_RGBA) ;`
- Double buffering must be specified during the context creation:
  - `glutInitDisplayMode(... | GLUT_DOUBLE) ;`
  - Once finished with the rendering of the current frame, front-/back-buffer swapping must be explicitly invoked:
    - `glutSwapBuffers() ;`




# Framebuffer

- Framebuffer clear color is specified through its RGBA components:
  - `glClearColor(red, green, blue, alpha);`
    - Alpha is required in RGB mode, too.
- Framebuffer is cleared explicitly through:
  - `glClear(GL_COLOR_BUFFER_BIT);`
- Details about the current context are retrieved through:
  - `glutGet(enum);`
    - E.g.: `GLUT_WINDOW_BUFFER_SIZE`, `GLUT_WINDOW_RED_SIZE`, `GLUT_WINDOW_GREEN_SIZE`, `GLUT_WINDOW_DOUBLEBUFFER`, ...



## Z buffer



- Is used to store the depth (z) value of each pixel of the framebuffer.
- Works like a “sonar”.
- Operations to the framebuffer are conditioned by the z buffer current state.






## Z buffer


- Without depth test:



a) first the triangle,  
then the rectangle





b) first the rectangle,  
then the triangle




## Z buffer


- With depth test: order-independent rendering




a) first the triangle,  
then the rectangle




b) first the rectangle,  
then the triangle



# Z buffer



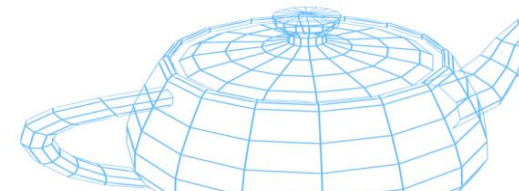
A simple three-dimensional scene



Z-buffer representation

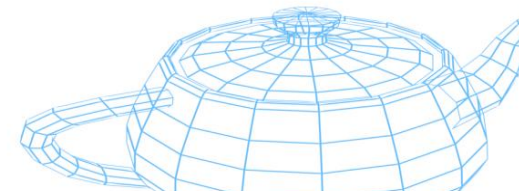
## Z buffer

- Contains the perpendicular distance of each pixel of the framebuffer relative to the near clipping plane.
- Values in normalized device coordinates  $[-1 < z < 1]$  are resampled into the  $[0 < z < 1]$  range.
- Accuracy depends on the number of bit used:
  - Typically 16/24 bit.
  - The 24 bit z buffer is often padded with an 8 bit stencil buffer to reach 32 bit boundaries.
  - 32 bit z buffers are possible only through off-screen framebuffers (via modern OpenGL framebuffer objects).




## Z buffer

- Precision is higher for objects closer to the *zNear* plane.
- In general, try to push the *zNear* plane out and the *zFar* plane in as much as possible:
  - Your z buffer accuracy corresponds to the discretization of the space  $zFar - zNear$  using  $X$  bit, where  $X$  is your z buffer bit depth.
- **Always** keep  $zNear > 0$ .



## Z buffer


- Must be specified during the context creation:
  - `glutInitDisplayMode(... | GLUT_DEPTH) ;`
- Must be explicitly enabled:
  - `glEnable(GL_DEPTH_TEST) ;`
- Z buffer must be cleaned before rendering:
  - `glClear(... | GL_DEPTH_BUFFER_BIT) ;`
- Z buffer behavior must be configured:
  - `glClearDepth(float) ;` [default is 1.0]
  - `glDepthFunc(enum) ;` [default is `GL_LESS`]




## Stencil buffer



- Optional buffer with the same dimension of the color and z buffers and a typical depth of 8 bit:
  - Could be just 1 bit.
  - Interaction with the z buffer:
    - Stencil buffer values can be modified according to the result of the depth test.
- Mainly used to limit the rendering to specific, pixel-precise areas:
  - Planar reflections.
- Other advanced applications involve volume shadows, constructive solid geometry, portals, etc.







## Stencil buffer

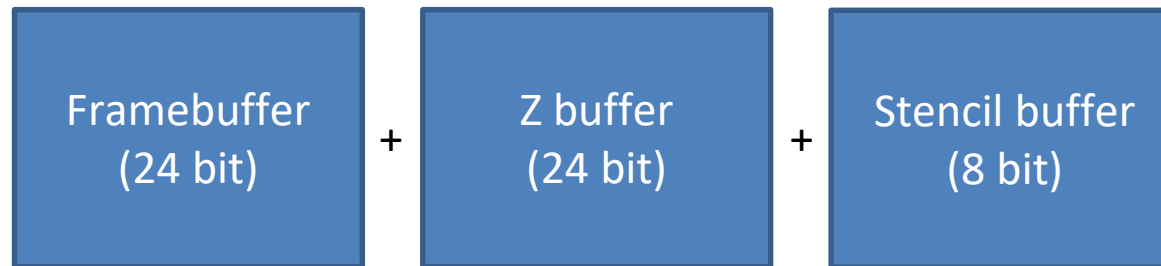


- Must be specified during the context creation:
  - `glutInitDisplayMode(... | GLUT_STENCIL);`
- Must be explicitly enabled:
  - `glEnable(GL_STENCIL_TEST);`
- Stencil buffer must be cleaned before rendering:
  - `glClear(... | GL_STENCIL_BUFFER_BIT);`
- Stencil buffer's behavior must be configured:
  - `glClearStencil(int);` [default is 0]
  - `glStencilFunc(enum, ref, mask);`
  - `glStencilOp(fail, zfail, zpass);`



## Main buffers

OpenGL context 1 (640x480, no double buffer):




640x480 = 307'200  
pixels  
24+24+8 = 56 bit  
~2 MB of VRAM  
(48 MB/s @ 24 fps)

OpenGL context 2 (1920x1080, double buffer):

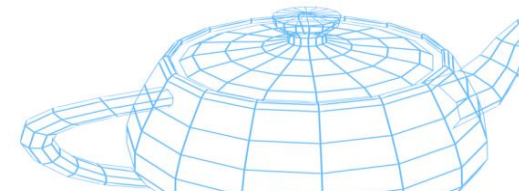


1920x1080 = 2'073'600 pixels  
32+24 = 56 bit  
~14 MB of VRAM (~20 MB x2)  
(336 MB/s @ 24 fps)



## Per-vertex information


- Vertex position (as seen so far):
  - $x, y, z, w$  (usually as *float*)
- Vertex color (RGB or RGBA):
  - $r, g, b, a$  (usually as *byte*)
- ...we will see additional per-vertex data later in the course (like normal vectors and texture coordinates).



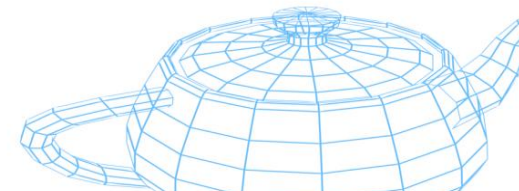
## Immediate mode




- Generates primitives according to the type specified and the number of vertices passed.
- A new vertex is generated when `glVertex*()` is called, using the last color values specified.



```
glBegin(GL_TRIANGLES);  
    glColor3f(1.0f, 0.0f, 0.0f);  
    glVertex3f(0.0f, 0.0f, 0.0f);  
    glColor3f(0.0f, 1.0f, 0.0f);  
    glVertex3f(10.0f, 0.0f, 0.0f);  
    glVertex3f(5.0f, 5.0f, 0.0f);  
glEnd();
```





World  
Transformations

```
glMatrixMode(GL_MODELVIEW);
glLoadMatrixf(m);
    (where m is typically cameraMat-1 * transMat * rotMat * scaleMat)
glBegin(...)
    glVertex* (...)
glEnd();
```

Lighting

Projection  
Transformations

```
glMatrixMode(GL_PROJECTION);
glLoadMatrixf(m);
    (where m is typically computed by glm::ortho or glm::perspective)
```

Clipping

```
glViewport(startX, startY, width, height);
```

Rasterization



$$\begin{bmatrix} clip_x \\ clip_y \\ clip_z \\ clip_w \end{bmatrix} = \overbrace{\text{projMat}}^{\text{GL\_PROJECTION}} * \overbrace{\text{cameraMat}^{-1} * \text{transMat} * \text{rotMat} * \text{scaleMat}}^{\text{GL\_MODELVIEW}} * \text{glVertex*}() \begin{bmatrix} obj_x \\ obj_y \\ obj_z \\ 1 \end{bmatrix}$$



Each time you resize the window, don't forget to update the projection matrix and `glViewport()` values accordingly!






## Matrices

- OpenGL stores, for each mode, the current matrix.
- Current matrix is set by passing a `glm::mat4` pointer to the method `glLoadMatrixf(float *)`;

```
#include <glm/gtc/type_ptr.hpp>

glMatrixMode(GL_MODELVIEW);
glm::mat4 mv = cameraInv * translation * rotation;
glLoadMatrixf(glm::value_ptr(mv));

glMatrixMode(GL_PROJECTION);
glm::mat4 pj = glm::perspective(...)
glLoadMatrixf(glm::value_ptr(pj));
```



## Matrices


- For any given matrix used to position an object in world coordinates:

$$\mathbf{M} = \begin{bmatrix} 1 & 0 & 0 & x \\ 0 & 1 & 0 & y \\ 0 & 0 & 1 & z \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

**x      y      z      t**

where **x**, **y**, **z** and **t** are column vectors representing the object local:

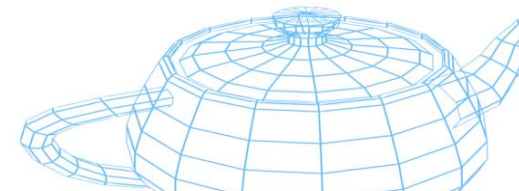
- Right direction (**x**).
- Up direction (**y**).
- Forward direction (**z**).
- Position (**t**).



## LookAt

- Commodity method for computing the camera matrix from a set of given parameters (eye, center and up):
  - **eye** (vec3): is the position of the camera.
  - **center** (vec3): is the position of the point the camera is looking at.
  - **up** (vec3): is a vector indicating the orientation of the world (typically 0, 1, 0).
- Available through the `glm::lookAt()` method:

```
glm::vec3 eye = glm::vec3(0.0f, 0.0f, 10.0f);  
glm::vec3 center = glm::vec3(0.0f, 0.0f, 0.0f);  
glm::vec3 up = glm::vec3(0.0f, 1.0f, 0.0f);  
  
glm::mat4 viewMat = glm::lookAt(eye, center, up);
```





# Camera demo

FOV: 54.0

Near plane: 1.0

Far plane: 50.0

Camera matrix (OpenGL order):

```
1.00, 0.00, 0.00, 0.00
0.00, 0.82, 0.57, 0.00
0.00, -0.57, 0.82, 0.00
0.00, -0.00, -8.84, 1.00
```

Projection matrix (OpenGL order):

```
2.41, 0.00, 0.00, 0.00
0.00, 1.96, 0.00, 0.00
0.00, 0.00, -1.04, -1.00
0.00, 0.00, -2.04, 0.00
```

Perspective =

$$= \begin{bmatrix} \frac{f}{\text{aspect}} & 0 & 0 & 0 \\ 0 & f & 0 & 0 \\ 0 & 0 & \frac{z_{\text{Far}} + z_{\text{Near}}}{z_{\text{Near}} - z_{\text{Far}}} & \frac{2 \times z_{\text{Far}} \times z_{\text{Near}}}{z_{\text{Near}} - z_{\text{Far}}} \\ 0 & 0 & -1 & 0 \end{bmatrix}$$



# Tutorial

Advanced FreeGLUT