

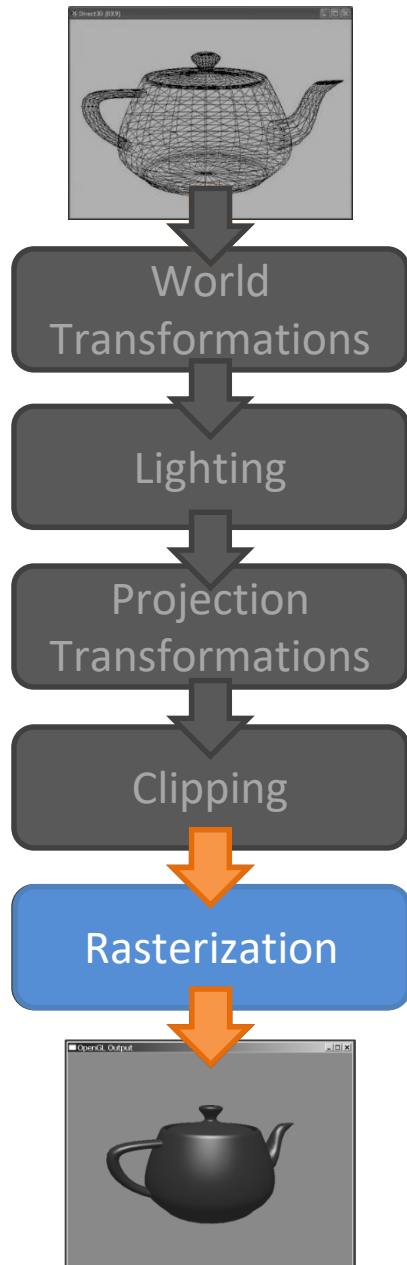
**SUPSI**

# Computer Graphics


Transparency

Achille Peternier, lecturer

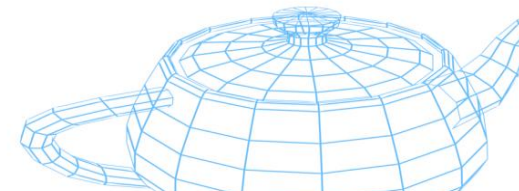




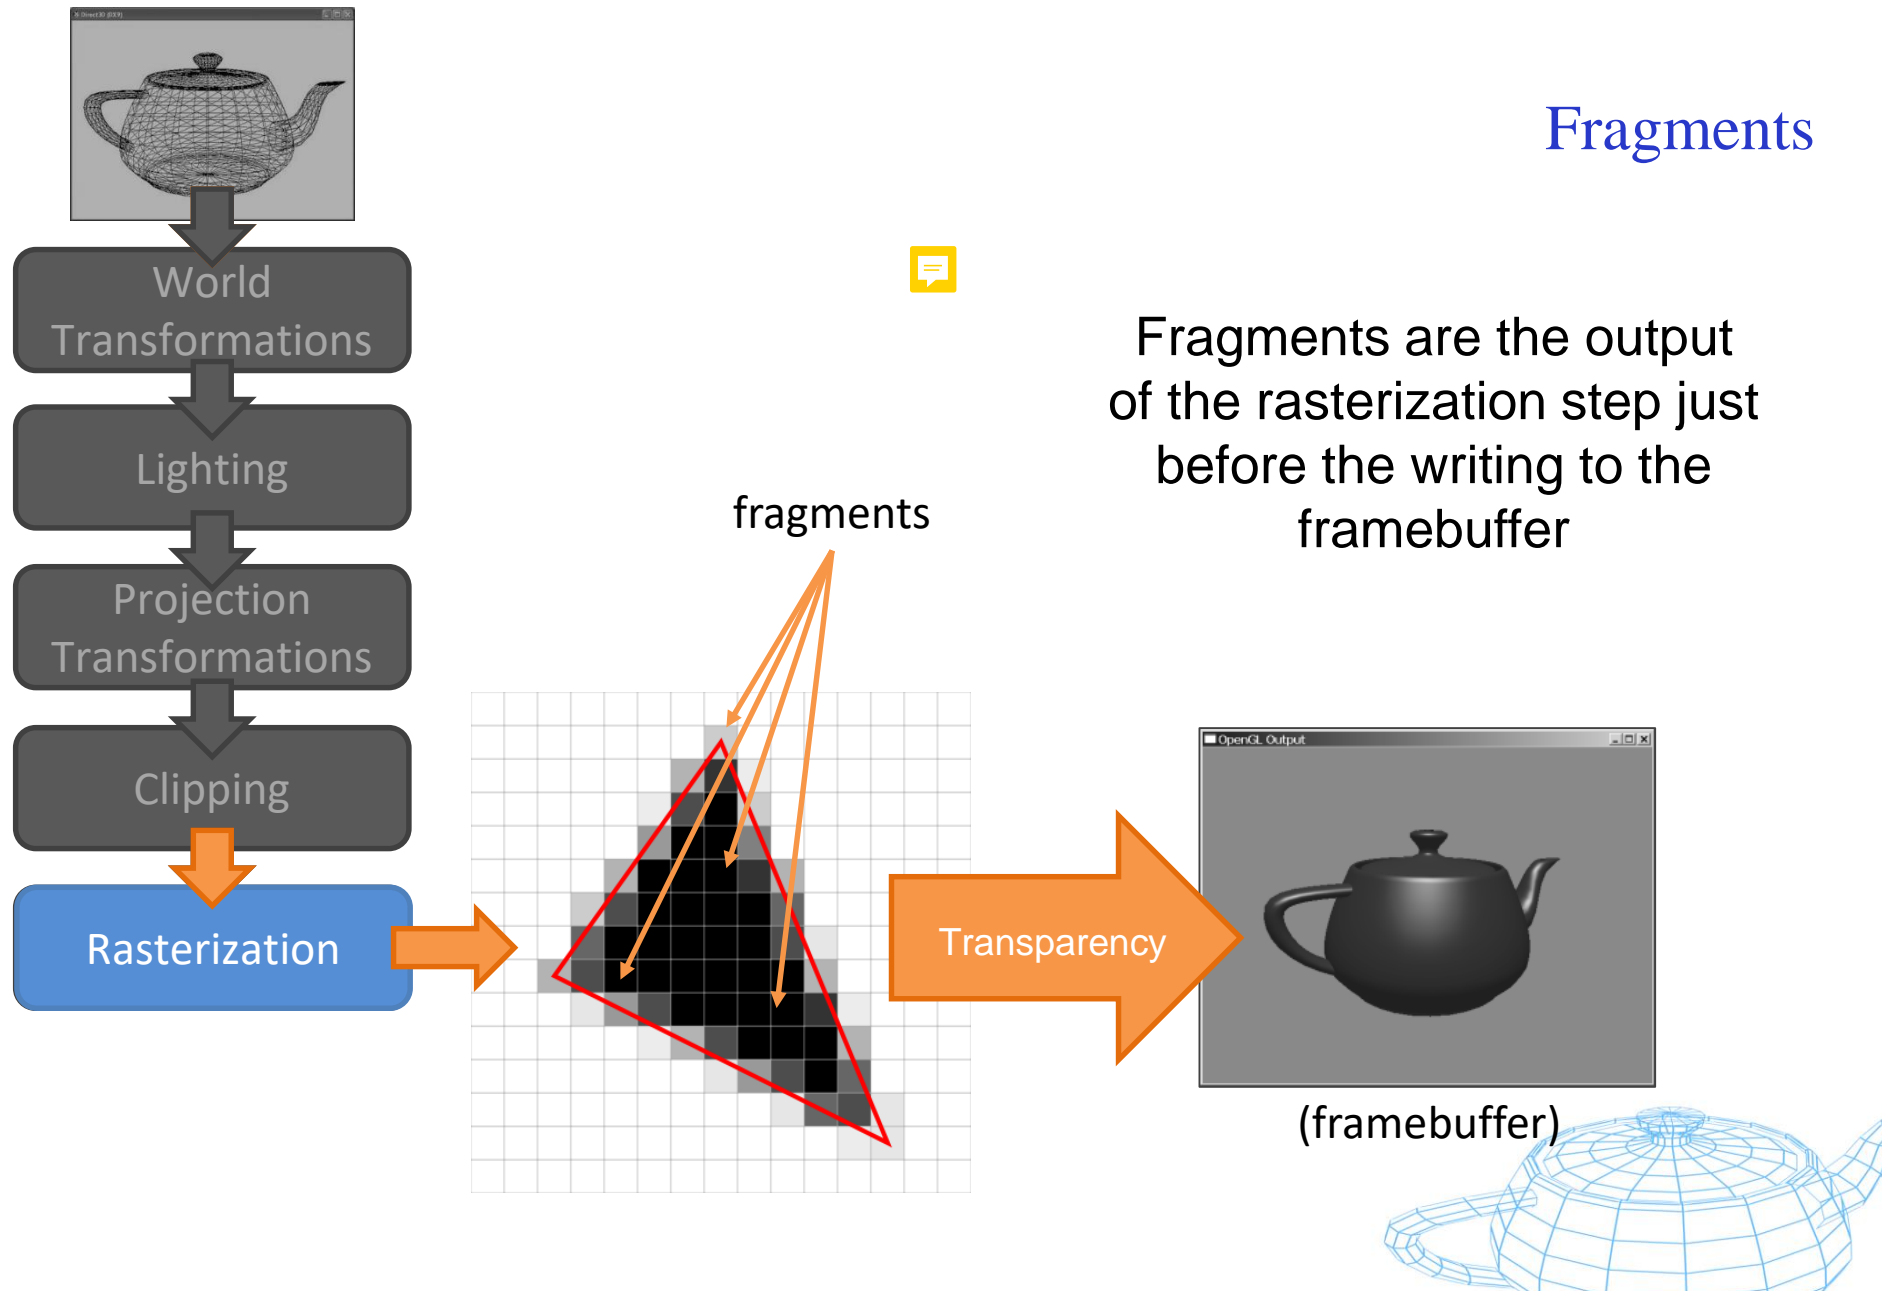
Transparent plane



Transparency



## Fragments



## Without transparency




- When transparency is not used, output fragments generated during the rasterization are copied to the framebuffer (as long as they pass the depth test).
- Solid/opaque fragments simply replace the different color pixels already stored in the framebuffer.




fragments


*New fragments that pass the z-test  
replace the framebuffer content*



No transparency




(framebuffer)




# Transparency

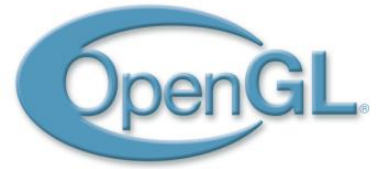
- There are two main types of transparency: **hard** and **soft**.
- **Hard transparency** refers to a binary kind of transparency, where an element is either entirely solid (not transparent at all) or completely transparent (thus invisible):
  - Hard transparency is implemented in OpenGL through the **alpha test**.
- **Soft transparency** enables progressive transparency with mutable gradients of intensity to provide see-through effects:
  - Soft transparency is supported in OpenGL through **blending**.
- Transparency intensity is commonly expressed through the alpha value (e.g., **RGBA**, `glColor4f()`, texture alpha channel, `glMaterial()` alpha component, etc.).



## Transparency

- In OpenGL, transparency is applied in the **blending** step:
  - Blending is the last operation of the pipeline and consists in copying one fragment (as an element of a rendered primitive) to the framebuffer.
  - When transparency is used, instead of simply overwriting pixels in the framebuffer with new content, fragments are merged (added, multiplied, interpolated, etc.) with the pixels already there.
- Blending is a computational expensive task: each fragment requires an additional framebuffer memory read and other operations to be done:
  - Did you ever notice the fps drop when smoke, volumetric fog, or explosions are rendered?






## Alpha test

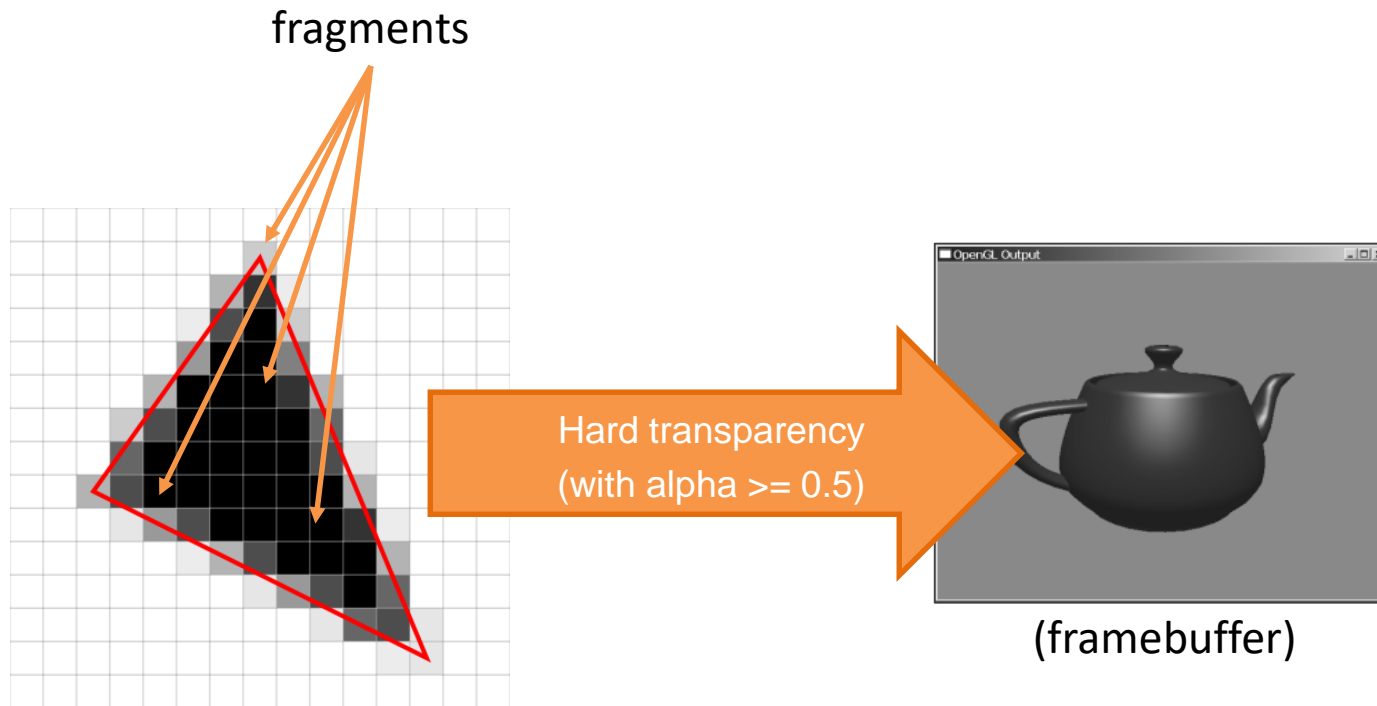
- **Hard transparency:** a pixel can be either solid or invisible:
  - Alpha test is activated through `glEnable(GL_ALPHA_TEST)`.
- A fragment is discarded when its alpha value is above/below a specific threshold:
  - Customize the reference threshold using the command `glAlphaFunc(function, value)`:
    - E.g.: `glAlphaFunc(GL_GREATER, 0.5f)` ;
- Sorting is not required when hard transparency is used:
  - The depth test works as usual.
- In modern OpenGL, alpha testing can be directly implemented in fragment shaders, e.g.:

*if (pixel.a < 0.5f) discard;*






## Alpha test




*New fragments that pass the z-test  
and with an alpha value of 0.5 or  
higher replace the framebuffer  
content*





# Examples



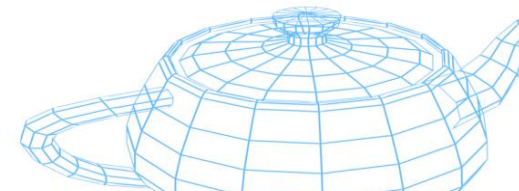
9





## Blending

- Blending is used to combine values already written in the framebuffer with the fragments added by a new primitive.
- Blending is activated through the command `glEnable(GL_BLEND)`.
- Different options are available to specify how pixels in the framebuffer interact with the new fragments:
  - Use `glBlendFunc(sourceFactor, destFactor)` to change OpenGL behavior.






## Blending

- Blending simply works as follows:

$$fbOutput = sourceFactor \times sourceColor + destFactor \times destColor$$

(where **fbOutput** is the color of the pixel written to the framebuffer, **sourceColor** is the fragment output pixel and **destColor** is the color of the pixel already in the framebuffer)

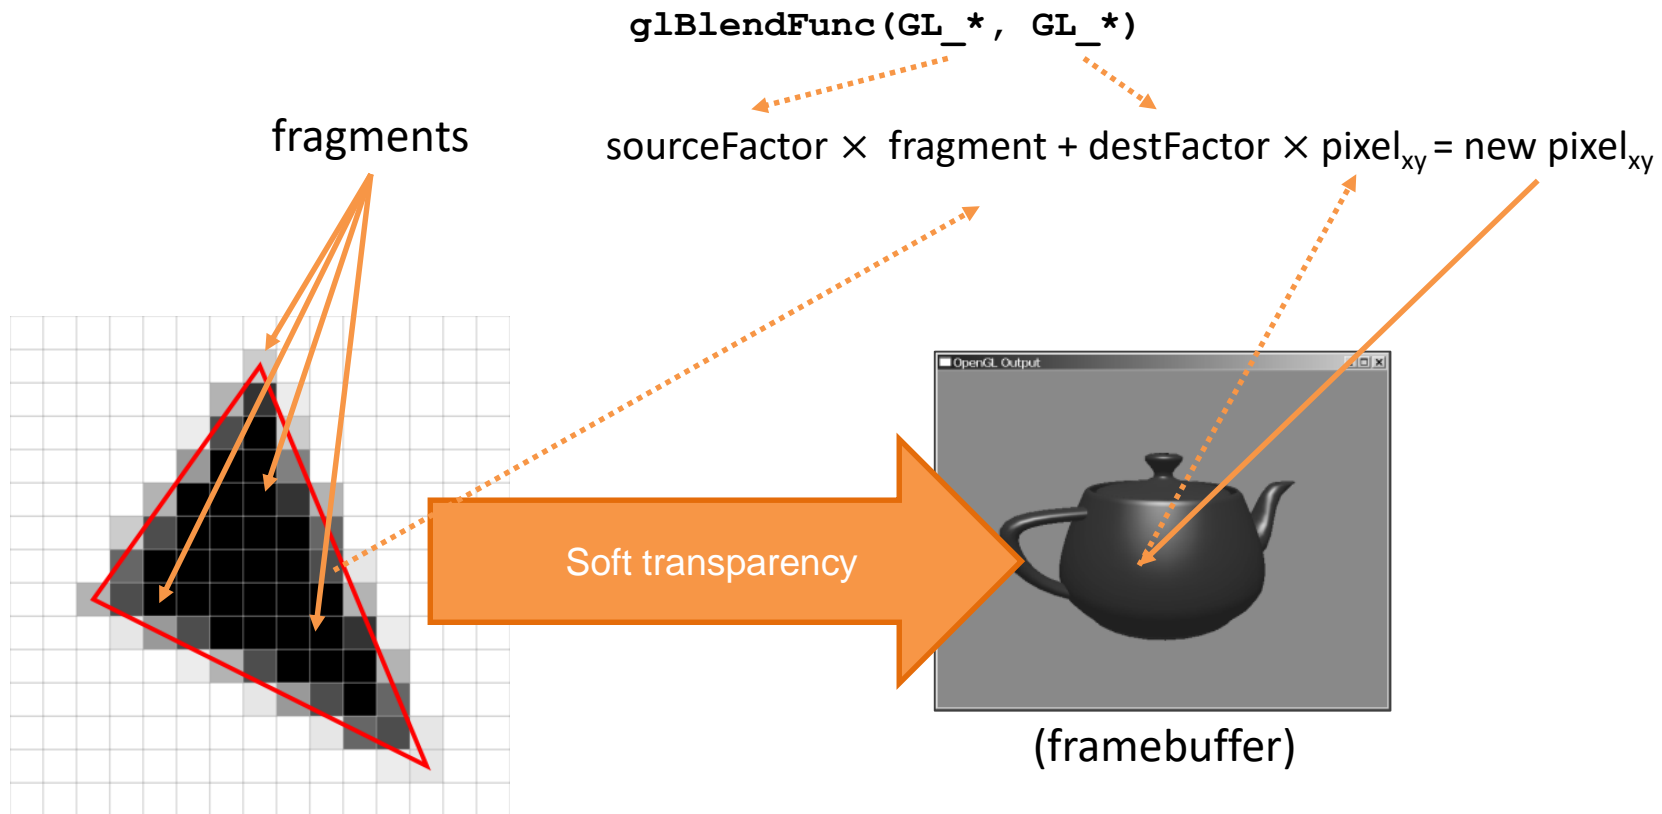
- Factors can be any combination of `GL_ZERO`, `GL_ONE`, `GL_DST_COLOR`, `GL_ONE_MINUS_DST_COLOR`, `GL_SRC_ALPHA`, `GL_ONE_MINUS_SRC_ALPHA`, `GL_DST_ALPHA`, `GL_ONE_MINUS_DST_ALPHA`, and `GL_SRC_ALPHA_SATURATE`.
- Typical transparency is implemented through the linear interpolation of two values using `glBlendFunc(GL_SRC_ALPHA, GL_ONE_MINUS_SRC_ALPHA)`.



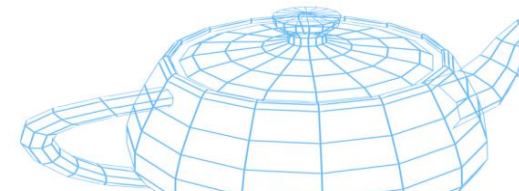




## Blending



*New fragments that pass the z-test replace the framebuffer content with a new pixel computed as the output of the blending function that takes the fragment color and the previous framebuffer pixel as input.*



# Particle engine demo

Exit module

Particle parameters

Energy: 80

Size: 100

System

Num

Spin

System

X: 0.0

Y: 0.0

Z: 0.0

Gravity

X: 0.00

Y: 23.3

Z: 0.00

System

X: 6.00

Y: -180

Z: 0.00

Blending options

Select texture...

ual Reality Laboratory

GL\_ONE\_MINUS\_SRC\_ALPHA

GL\_ONE\_MINUS\_SRC\_ALPHA

GL\_ONE\_MINUS\_DST\_ALPHA

GL\_ONE\_MINUS\_DST\_ALPHA


GL\_SRC\_ALPHA\_SATURATE

## Blending problems

- Most soft transparency blending operations are order-dependent:




a) Green first, then red:



RGB: [194 83 60]


WARNING

b) Red first, then green:



RGB: [102 147 89]


!=





## Blending problems



- Correct transparency works only when transparent triangles are rendered back to front:
  - Real-time sorting is computationally expensive on a per-triangle basis.
  - Interleaved triangles are not rendered properly.
- Mesh triangles are always passed in the same order (and are most likely rendered in the same sequence):
  - Depending on the mesh and camera orientation, their sorting can be correct or wrong.



## Multi-pass rendering

- **Single-pass rendering:** all the elements of the scene are rendered within one single loop/iteration:
  - Each object is rendered in its final state, including textures, multiple light contributions, etc. 
- **Multi-pass rendering:** the scene is generated after multiple iterations, where at each step additional information is provided:
  - E.g.: rendering the full scene once per each light source available, then use additive blending to merge the output images together to generate the final image. 

1

2

3

1 + 2 + 3

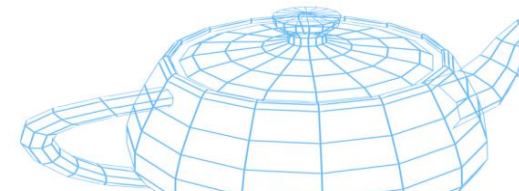




## Transparency rendering



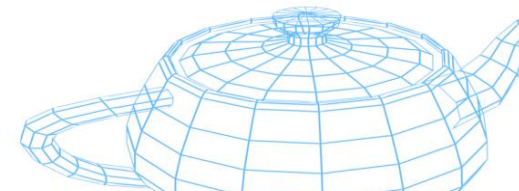
- Relatively artifact-free transparency rendering can be achieved through a multi-pass approach:
  - 1<sup>st</sup> pass:
    - Clear buffers, enable depth testing.
    - Render all the solid objects (as we did so far).
  - 2<sup>nd</sup> pass:
    - Set depth testing to read-only via `glDepthMask()`.
    - Depth-sort transparent objects (farthest to nearest).
    - Render each transparent object in two sub-passes:
      - Sub-pass 1: enable front-face culling.
      - Sub-pass 2: enable back-face culling.






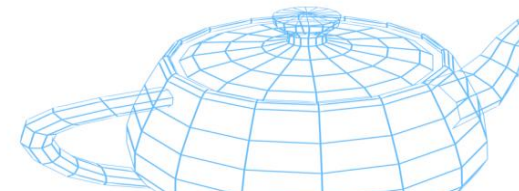
## Transparency rendering

- The previous technique can be easily implemented in your graphics engine's list class:
  - The list already contains all the objects to be rendered, including their world matrix.
  - Sort transparent objects according to distance between the positional part of their world matrix and the positional part of the camera matrix:
    - Make sure to compare positions defined in the same space, e.g., world coordinates or eye coordinates.
  - You can refine results and limit cross-penetrations by applying a bounding sphere to each mesh and by including its radius in the depth sorting:
    - The OVO format includes information about each mesh's bounding radius.




## Graphics engine remarks

- Class material: 
  - Add an **isTransparent()** method to quickly identify transparent materials.
  - Consider adding a **setTransparency()** method to change the alpha value of ambient, diffuse, and specular with one single instruction.
- Class list:
  - Implement a depth-sorting (back to front) method for transparent meshes.
- Class mesh:
  - If a mesh uses a transparent material, render it in two passes using different face culling parameters during each pass.
- Improve font readability by writing on top of semitransparent rectangles, e.g.:




## Transparency rendering

- The previous technique is a compromise between implementation complexity, rendering quality and computational power required:
  - It requires depth-sorting:
    - Although not on a per-triangle basis.
  - It does not work on nested or interpenetrating meshes.
  - It does not work on meshes with complex shapes, e.g.:



- Several Order-Independent Transparency (OIT) techniques exist (such as A-buffers and depth peeling) but require a recent version of OpenGL and additional knowledge.



## Depth peeling (OIT technique)

- The idea is to use a more complex framebuffer with multiple layers (each one with its own depth buffer).
- Multi-pass rendering is used (one pass per layer):
  - The scene is rendered.
  - Only the nearest pixels are stored in the first layer (as we did so far).
  - Rendering is then done again, once for each layer:
    - At each iteration, the 2<sup>nd</sup>, 3<sup>rd</sup>, ..., n<sup>th</sup> nearest pixels are stored within each layer.
- Once all the layers have been processed, they are rendered back to front.

