

SUPSI

Programming CUDA

Tiziano Leidi

Sviluppo di Applicazioni High-Performance
Bachelor in Ingegneria Informatica

C for CUDA

CUDA provides several **additions to the C/C++ environment:**

- C language extensions to enable heterogeneous programming
- mathematical functions of the C/C++ standard library (for both host and device)
- intrinsic functions (only supported on the device)
- CUDA runtime API and CUDA driver API
- Environment variables

C/C++ for the device code has some limitations compared to the full standard. For example it does not support run time type information (RTTI) and exception handling.

Example: Device Code

A simple kernel to add two vectors of integers:

```
__global__ void add(int *a, int *b, int *c) {  
    int index = threadIdx.x + blockIdx.x * blockDim.x;  
    c[index] = a[index] + b[index];  
}
```

CUDA Kernel

- The function type qualifier `__global__` declares a function as being an executable kernel on the CUDA device.
- This function can only be called from the host.
- All kernels **must be declared with a return type of void.**

Example: Host Code

```
#define N (2048*2048)
#define THREADS_PER_BLOCK 512
int main(void) {
    int *a, *b, *c;           // host copies of a, b, c
    int *d_a, *d_b, *d_c;     // device copies of a, b, c
    int size = N * sizeof(int);

    // Alloc space for device copies of a, b, c
    cudaMalloc((void **)&d_a, size);
    cudaMalloc((void **)&d_b, size);
    cudaMalloc((void **)&d_c, size);

    // Alloc space for host copies of a, b, c
    // and setup input values
    a = (int *)malloc(size); random_ints(a, N);
    b = (int *)malloc(size); random_ints(b, N);
    c = (int *)malloc(size);
```

Example: Host Code

```
// Copy inputs to device
```

```
cudaMemcpy(d_a, a, size, cudaMemcpyHostToDevice);
```

```
cudaMemcpy(d_b, b, size, cudaMemcpyHostToDevice);
```

```
// Launch add() kernel on GPU
```

```
add<<<N/THREADS_PER_BLOCK, THREADS_PER_BLOCK>>>(d_a, d_b, d_c);
```

```
// Copy result back to host
```

```
cudaMemcpy(c, d_c, size, cudaMemcpyDeviceToHost);
```

```
// Cleanup
```

```
free(a); free(b); free(c);
```

```
cudaFree(d_a); cudaFree(d_b); cudaFree(d_c);
```

```
return 0;
```

```
}
```

Function Type Qualifiers

The available function type qualifiers are:

__device__

- Executed on the device
- Callable from the device only

__global__

- Executed on the device
- Callable from the host only

__host__

- Executed on the host
- Callable from the host only

Calling the Kernel

- The host calls the kernel by specifying the name of the kernel plus an **execution configuration**:

```
Func<<< Dg, Db, Ns >>>(parameter) ;
```

- The configuration includes three things:
 - The number of threads in a group - **the block geometry**
 - The number of groups - **the grid geometry**
 - The number of bytes in shared memory that is dynamically allocated (optional) in addition to the statically allocated memory

dim3

- To pass the grid and block dimensions in a kernel invocation an integer vector type is used: **dim3**.
- dim3 has 3 elements x, y and z.
- In C code, dim3 can be initialized as:

```
dim3 grid = { 512, 512, 1 };
```
- In C++ code, dim3 can be initialized as:

```
dim3 grid( 512, 512, 1 );
```
- Not all the 3 elements need to be provided. Any element not provided during initialization is initialized to 1.

Calling the Kernel

Dimensions:

- *Dg* is of type **dim3**:
 $Dg.x * Dg.y$ = number of blocks being launched;
- *Db* is of type **dim3**:
 $Db.x * Db.y * Db.z$ = number of threads per block;
- *Ns* is of type **size_t**.

Built-in variables:

- *gridDim* is of type **dim3** - dimensions of the grid.
- *blockIdx* is of type **uint3** - block index within the grid.
- *blockDim* is of type **dim3** - dimensions of the block.
- *threadIdx* is of type **uint3** - thread index within the block.

Threads and Blocks

Thread 0

`c[0] = a[0] + b[0];`


Thread 1

`c[1] = a[1] + b[1];`

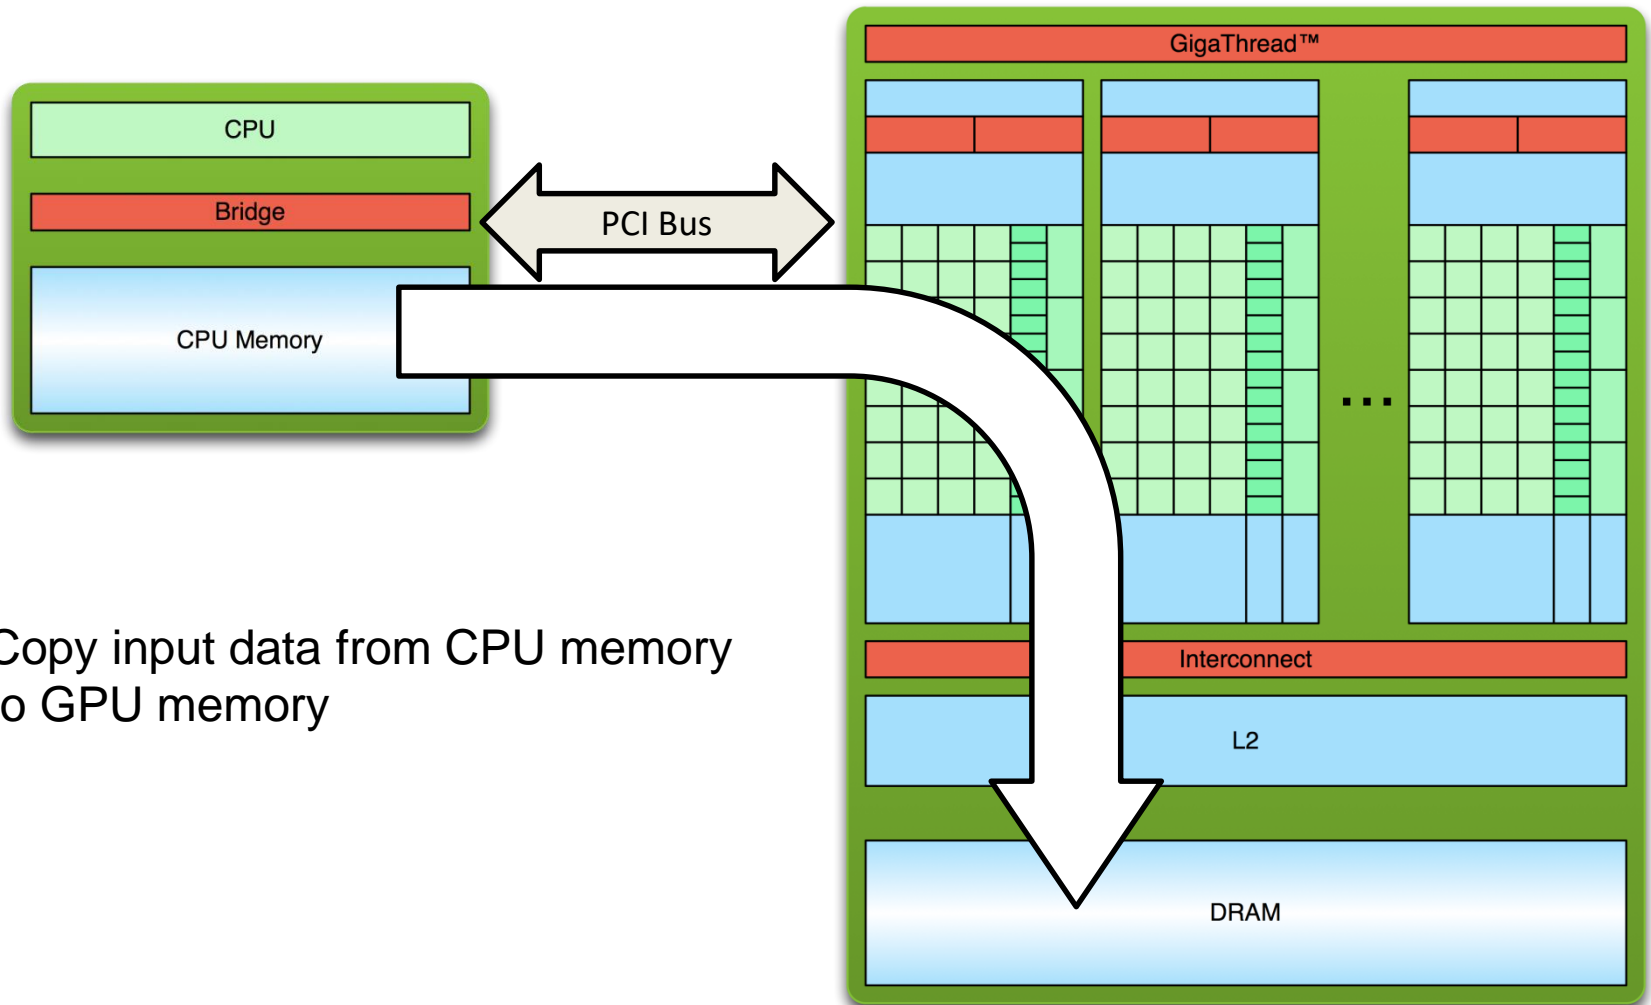
Thread 2

`c[2] = a[2] + b[2];`

Thread 3

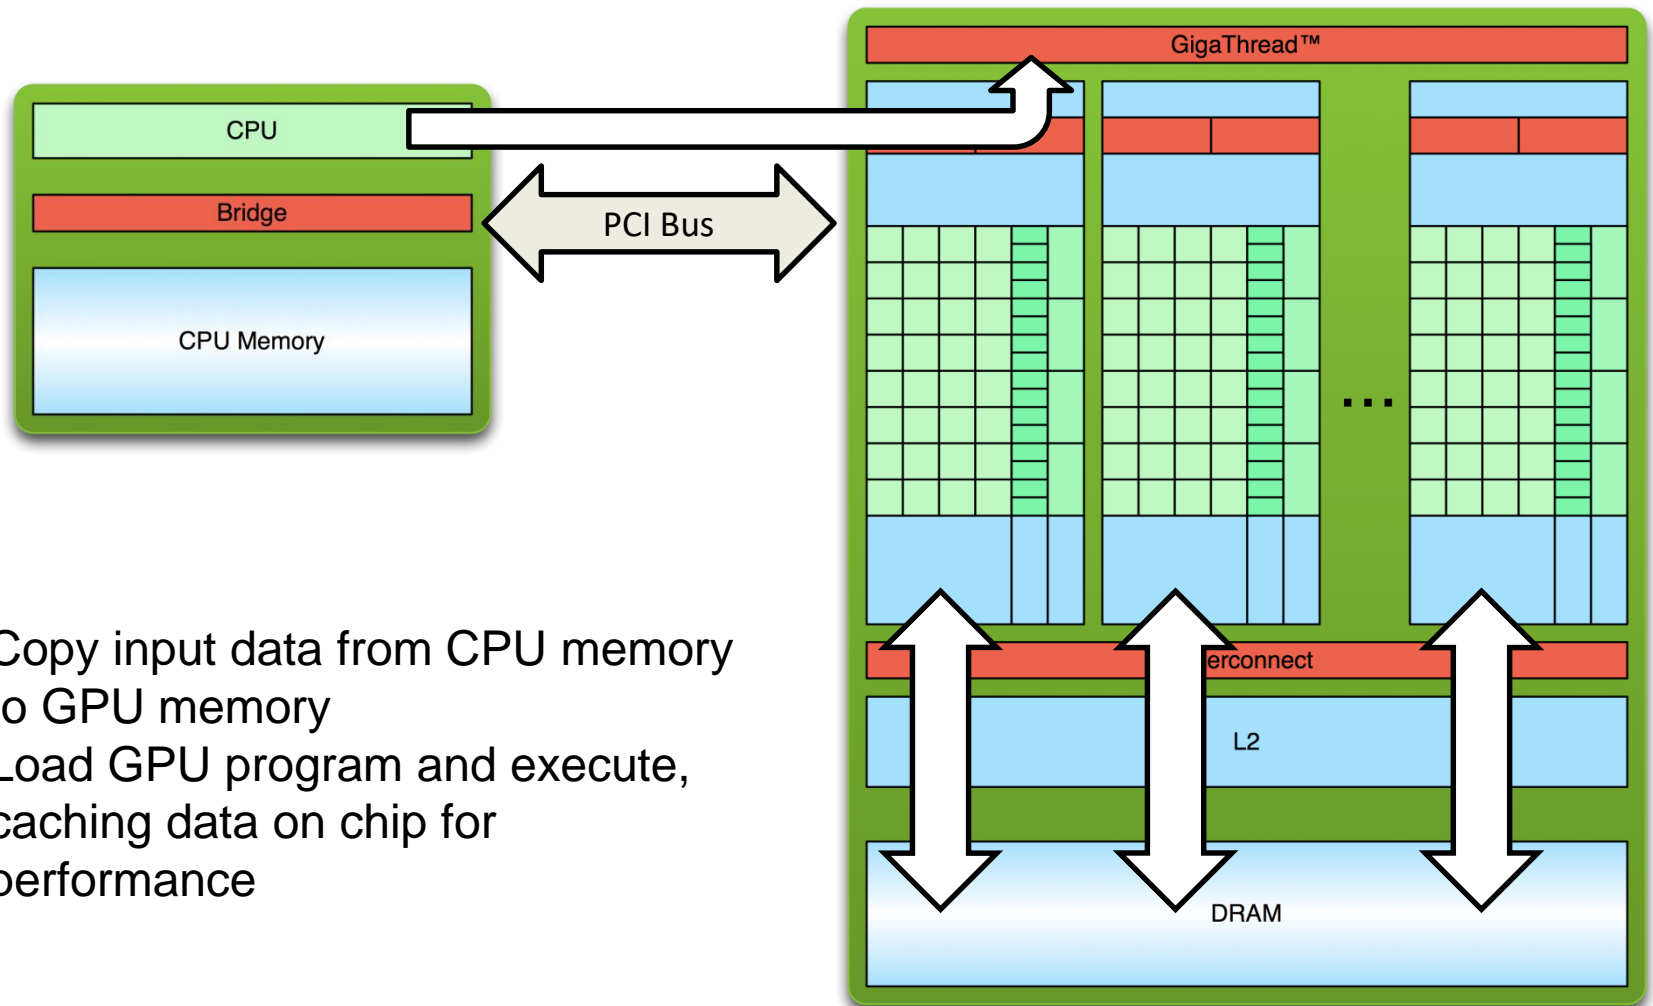
`c[3] = a[3] + b[3];`

Processing Flow

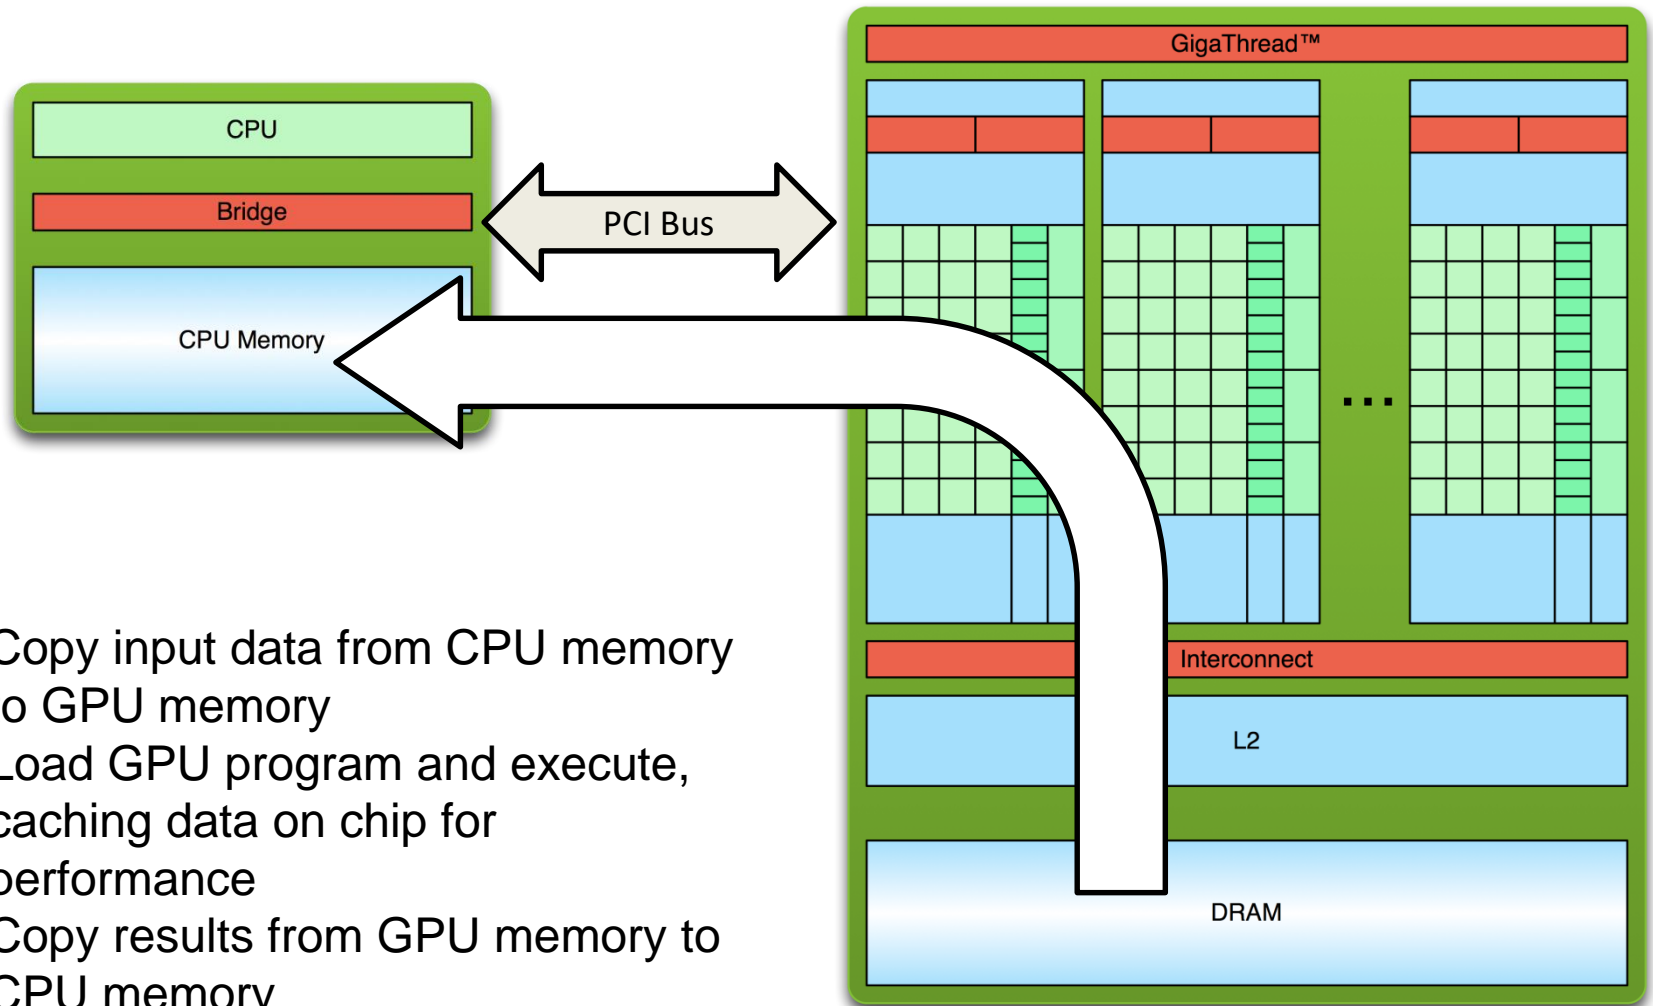


1. Copy input data from CPU memory to GPU memory

Processing Flow



Processing Flow

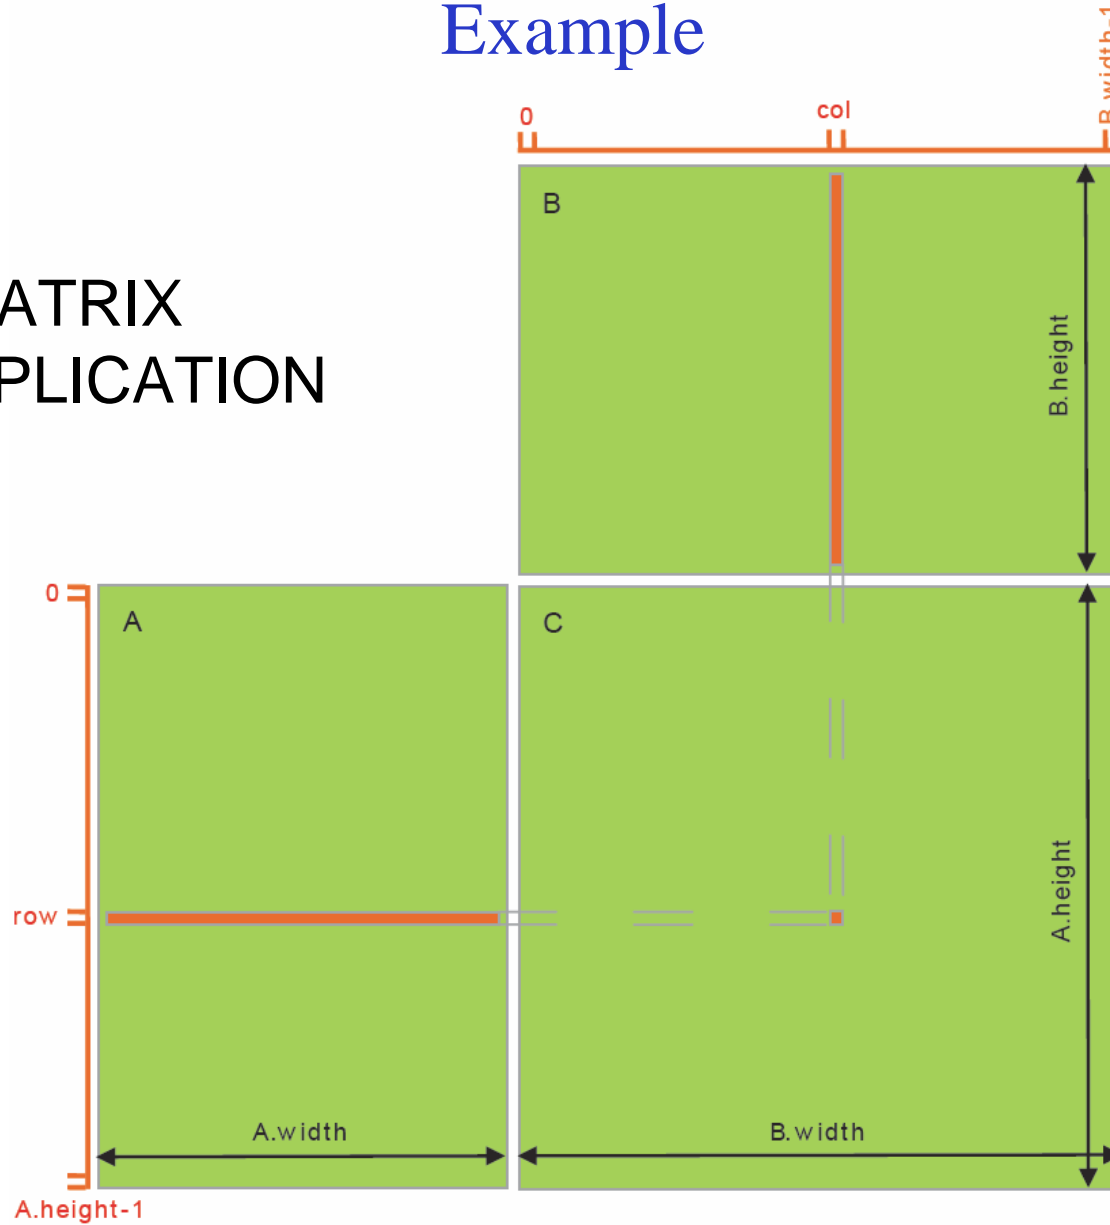


Compilation

- CUDA C language programs have **the suffix ".cu"**.
- Can be compiled directly to executable or to object code.
- Can be later linked to objects created by different compilers (e.g. integrated into visual studio).

Example

MATRIX MULTIPLICATION



Example

```
// Matrices are stored in row-major order:
// M(row, col) = *(M.elements + row * M.width + col)
typedef struct {
    int width;
    int height;
    float* elements;
} Matrix;

// Thread block size
#define BLOCK_SIZE 16

// Forward declaration of the matrix multiplication kernel
__global__ void MatMulKernel(const Matrix, const Matrix, Matrix);

// Matrix multiplication - Host code
// Matrix dimensions are assumed to be multiples of BLOCK_SIZE
void MatMul(const Matrix A, const Matrix B, Matrix C) {
```

Example

// Load A and B to device memory

Matrix d_A;

d_A.width = A.width; d_A.height = A.height;

size_t size = A.width * A.height * sizeof(float);

cudaMalloc(&d_A.elements, size);

cudaMemcpy(d_A.elements, A.elements, size, cudaMemcpyHostToDevice);

Matrix d_B;

d_B.width = B.width; d_B.height = B.height;

size = B.width * B.height * sizeof(float);

cudaMalloc(&d_B.elements, size);

cudaMemcpy(d_B.elements, B.elements, size, cudaMemcpyHostToDevice);

// Allocate C in device memory

Matrix d_C;

d_C.width = C.width; d_C.height = C.height;

size = C.width * C.height * sizeof(float);

cudaMalloc(&d_C.elements, size);

Example

```
// Invoke kernel
dim3 dimBlock(BLOCK_SIZE, BLOCK_SIZE);
dim3 dimGrid(B.width / dimBlock.x, A.height / dimBlock.y);
MatMulKernel<<<dimGrid, dimBlock>>>(d_A, d_B, d_C);
// Read C from device memory
cudaMemcpy(C.elements, Cd.elements, size, cudaMemcpyDeviceToHost);
// Free device memory
cudaFree(d_A.elements);
cudaFree(d_B.elements);
cudaFree(d_C.elements);
}
```

Example

```
// Matrix multiplication kernel called by MatMul()
__global__ void MatMulKernel(Matrix A, Matrix B, Matrix C) {
    // Each thread computes one element of C
    // by accumulating results into Cvalue
    float Cvalue = 0;
    int row = blockIdx.y * blockDim.y + threadIdx.y;
    int col = blockIdx.x * blockDim.x + threadIdx.x;
    for (int e = 0; e < A.width; ++e)
        Cvalue += A.elements[row * A.width + e] * B.elements[e * B.width + col];
    C.elements[row * C.width + col] = Cvalue;
}
```

Memory Management

CUDA provides a simple API for handling device memory:

- `cudaMalloc()`, `cudaFree()`, `cudaMemcpy()`
- similar to the C equivalents `malloc()`, `free()`, `memcpy()`

Host and device memory are separate entities.

Device pointers point to GPU memory:

- May be passed to/from host code
- May not be dereferenced in host code

Host pointers point to CPU memory:

- May be passed to/from device code
- May not be dereferenced in device code

Variable Type Qualifiers

The available variable type qualifiers are:

__device__

- Resides in **global memory space**
- Has the **lifetime of an application**
- Is accessible **from all the threads** within the grid and from the host through the runtime library

__constant__

- Resides in **constant (global) memory space**
- Has the **lifetime of an application**
- Is accessible **from all the threads** within the grid and from the host through the runtime library

Variable Type Qualifiers

__shared__

- Resides in the **shared memory space** of a thread block
- Has the **lifetime of the block**
- Is only accessible **from all the threads within the block**

Register Memory

Registers are:

- The **fastest form of memory** on the GPU
- Reside on the multi-core
- **Only accessible by the thread**
- Have the **lifetime of the thread**

Shared Memory

Shared memory:

- **Can be as fast as a registers** when there are no bank conflicts or when reading from the same address
- **Accessible by any thread of the block** from which it was created
- Has the **lifetime of the block**

Global Memory

Global memory:

- Potentially **150x slower** than register or shared memory
- Watch out for uncoalesced reads and writes
- Accessible from either the **host or device**
- Has the **lifetime of the application**

Local Memory

Local memory:

- A potential performance bottleneck
- Resides in global memory
- Can be 150x slower than register or shared memory
- Only accessible by the thread
- Has the lifetime of the thread

Used automatically by CUDA for:

- register overflow
- arrays addressed by parameters in register memory
- possibly misaligned global memory reads

Constant Memory

Constant memory:

- Cached access to global memory
- Read-only
- Limited to 64kb for the application
- Shared by all threads
- Shared by instruction cache
- Cache is 8kb

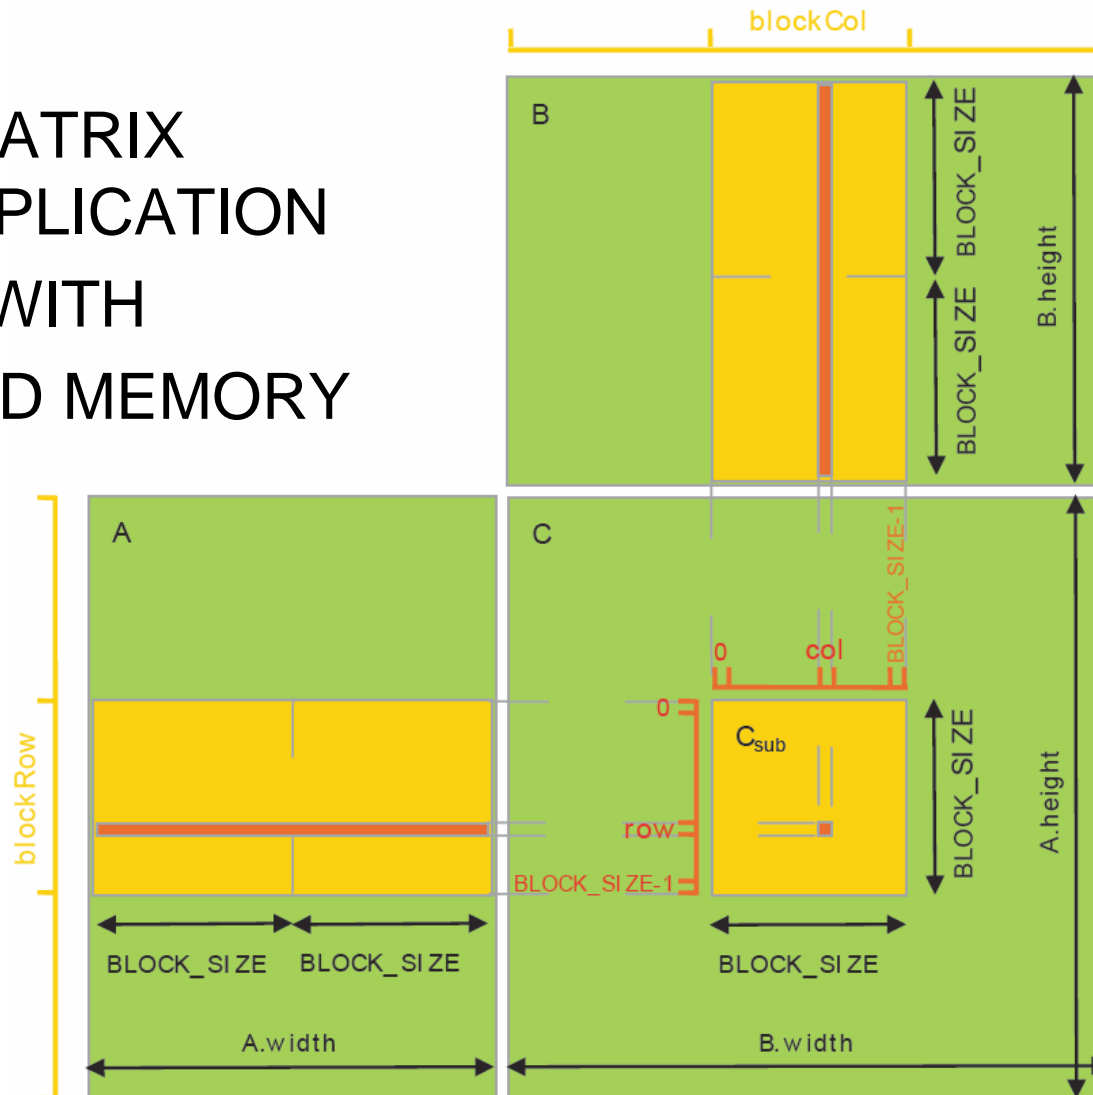
Texture Memory

Texture memory:

- Cached access to global memory
- Read-only
- L1 shared by multiple processors (depending on generation), L2 by all threads
- Optimized for 2D access
- L1 cache is 6kb to 8kb (depending on processor)

Example

MATRIX MULTIPLICATION WITH SHARED MEMORY



$B.width = stride$

Example

// Matrices are stored in row-major order:

// $M(\text{row}, \text{col}) = *(M.\text{elements} + \text{row} * M.\text{stride} + \text{col})$

typedef struct {

 int width;

 int height;

 int stride;

 float* elements;

*Larghezza della matrice di partenza,
necessaria per andare a capo
anche nel caso delle sottomatrici*

} Matrix;

// Get a matrix element

__device__ float GetElement(const Matrix A, int row, int col) {

 return A.elements[row * A.stride + col];

}

// Set a matrix element

__device__ void SetElement(Matrix A, int row, int col, float value) {

 A.elements[row * A.stride + col] = value;

}

Example

```
// Get the BLOCK_SIZExBLOCK_SIZE sub-matrix Asub of A that is  
// located col sub-matrices to the right and row sub-matrices down  
// from the upper-left corner of A
```

```
__device__ Matrix GetSubMatrix(Matrix A, int row, int col) {  
    Matrix Asub;  
    Asub.width = BLOCK_SIZE;  
    Asub.height = BLOCK_SIZE;  
    Asub.stride = A.stride;  
    Asub.elements = &A.elements[A.stride * BLOCK_SIZE * row + BLOCK_SIZE * col];  
    return Asub;  
}
```

```
// Thread block size
```

```
#define BLOCK_SIZE 16
```

```
// Forward declaration of the matrix multiplication kernel
```

```
__global__ void MatMulKernel(const Matrix, const Matrix, Matrix);
```


Example

```
// Matrix multiplication - Host code
// Matrix dimensions are assumed to be multiples of BLOCK_SIZE
void MatMul(const Matrix A, const Matrix B, Matrix C) {
    // Load A and B to device memory
    Matrix d_A;
    d_A.width = d_A.stride = A.width; d_A.height = A.height;
    size_t size = A.width * A.height * sizeof(float);
    cudaMalloc(&d_A.elements, size);
    cudaMemcpy(d_A.elements, A.elements, size, cudaMemcpyHostToDevice);
    Matrix d_B;
    d_B.width = d_B.stride = B.width; d_B.height = B.height;
    size = B.width * B.height * sizeof(float);
    cudaMalloc(&d_B.elements, size);
    cudaMemcpy(d_B.elements, B.elements, size, cudaMemcpyHostToDevice);
```

Example

```
// Allocate C in device memory
Matrix d_C;
d_C.width = d_C.stride = C.width; d_C.height = C.height;
size = C.width * C.height * sizeof(float);
cudaMalloc(&d_C.elements, size);
// Invoke kernel
dim3 dimBlock(BLOCK_SIZE, BLOCK_SIZE);
dim3 dimGrid(B.width / dimBlock.x, A.height / dimBlock.y);
MatMulKernel<<<dimGrid, dimBlock>>>(d_A, d_B, d_C);
// Read C from device memory
cudaMemcpy(C.elements, d_C.elements, size, cudaMemcpyDeviceToHost);
// Free device memory
cudaFree(d_A.elements);
cudaFree(d_B.elements);
cudaFree(d_C.elements);
}
```

Example

```
// Matrix multiplication kernel called by MatMul()
__global__ void MatMulKernel(Matrix A, Matrix B, Matrix C) {
    // Block row and column
    int blockRow = blockIdx.y;
    int blockCol = blockIdx.x;
    // Each thread block computes one sub-matrix Csub of C
    Matrix Csub = GetSubMatrix(C, blockRow, blockCol);
    // Each thread computes one element of Csub
    // by accumulating results into Cvalue
    float Cvalue = 0;
    // Thread row and column within Csub
    int row = threadIdx.y;
    int col = threadIdx.x;
```

Example

```
// Loop over all the sub-matrices of A and B that are required to compute Csub
// Multiply each pair of sub-matrices together and accumulate the results
for (int m = 0; m < (A.width / BLOCK_SIZE); ++m) {
    // Get sub-matrix Asub of A
    Matrix Asub = GetSubMatrix(A, blockRow, m);
    // Get sub-matrix Bsub of B
    Matrix Bsub = GetSubMatrix(B, m, blockCol);
    // Shared memory used to store Asub and Bsub respectively
    __shared__ float As[BLOCK_SIZE][BLOCK_SIZE];
    __shared__ float Bs[BLOCK_SIZE][BLOCK_SIZE];
    // Load Asub and Bsub from device memory to shared memory
    // Each thread loads one element of each sub-matrix
    As[row][col] = GetElement(Asub, row, col);
    Bs[row][col] = GetElement(Bsub, row, col);
    // Synchronize to make sure the sub-matrices are loaded
    // before starting the computation
    __syncthreads();
```

Example

```
// Multiply Asub and Bsub together
for (int e = 0; e < BLOCK_SIZE; ++e)
    Cvalue += As[row][e] * Bs[e][col];
    // Synchronize to make sure that the preceding
    // computation is done before loading two new
    // sub-matrices of A and B in the next iteration
    __syncthreads();
}
// Write Csub to device memory
// Each thread writes one element
SetElement(Csub, row, col, Cvalue);
}
```

Syncthreads

- **Threads within a block can cooperate** by sharing data through some shared memory and by synchronizing their execution to coordinate memory accesses.
- Synchronization points are specified in the kernel **by calling `__syncthreads()`**: a barrier at which all threads in the block must wait before any is allowed to proceed.
- For efficient cooperation, the shared memory is expected to be a low-latency memory near each processor core and **`__syncthreads()` is expected to be lightweight.**