

**SUPSI**

# Lab: The Bthread library

Operating Systems

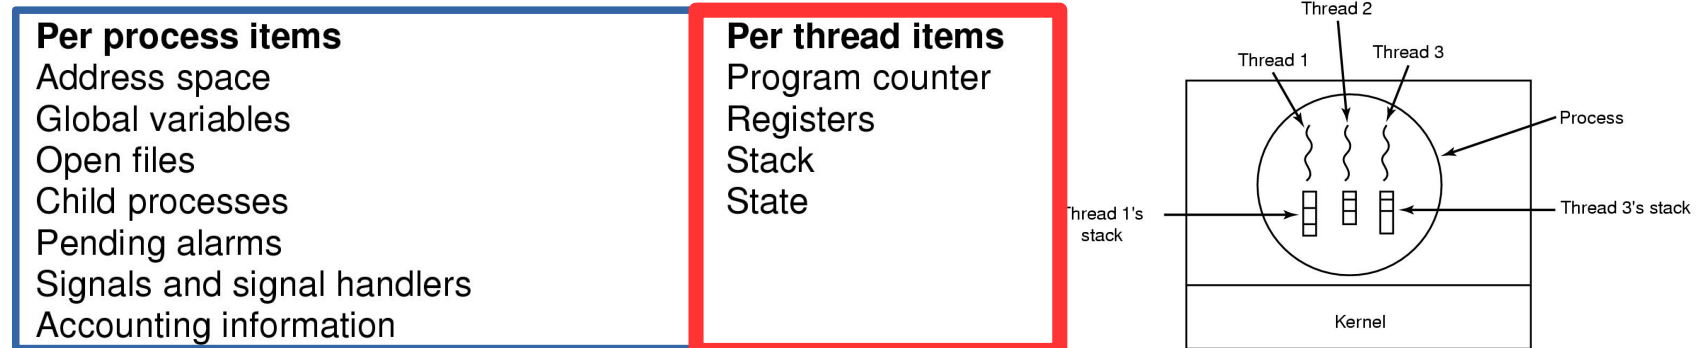
Amos Brocco, Lecturer & Researcher

## Objectives

- Understand the goal of the Bthread library
- Discover how to implement user-level threads

## The active part of a process: threads

- A process can have one or more **threads** (or **paths**) of execution \*
- Threads in a process share some resources (→ concurrency problems)

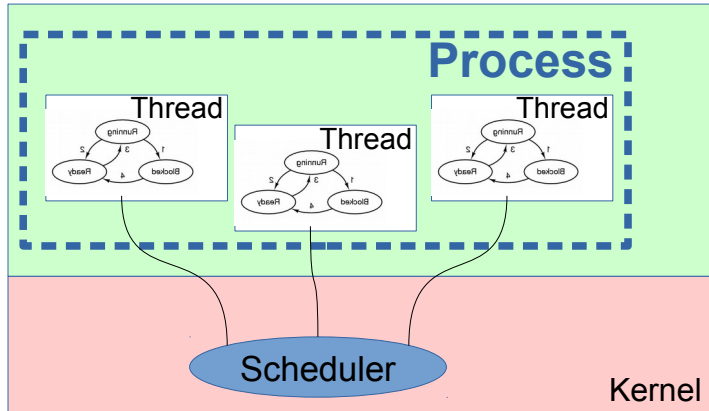


A. Tanenbaum, Modern Operating Systems, 2nd ed

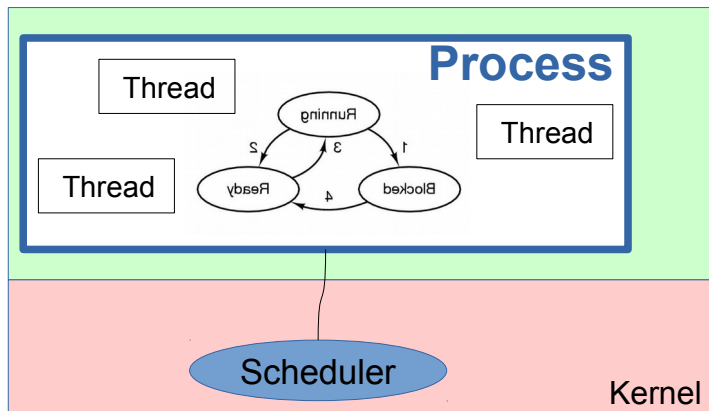
- When a process has multiple threads of execution we call it a **multi-threaded process**, otherwise it is called a **single-threaded process**

\* typically simply referred to as **threads**

# Threads implementation



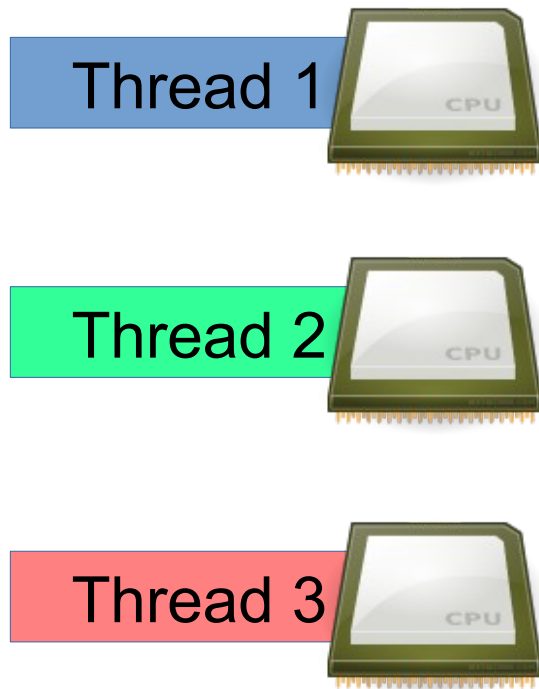
- **Kernel level threads**
  - Thread scheduling is done by the kernel
  - If a thread blocks, other threads within the same process can continue executing



- **User level threads**
  - Thread scheduling is done by the process
  - If a thread blocks, the whole process (including other user threads) is blocked

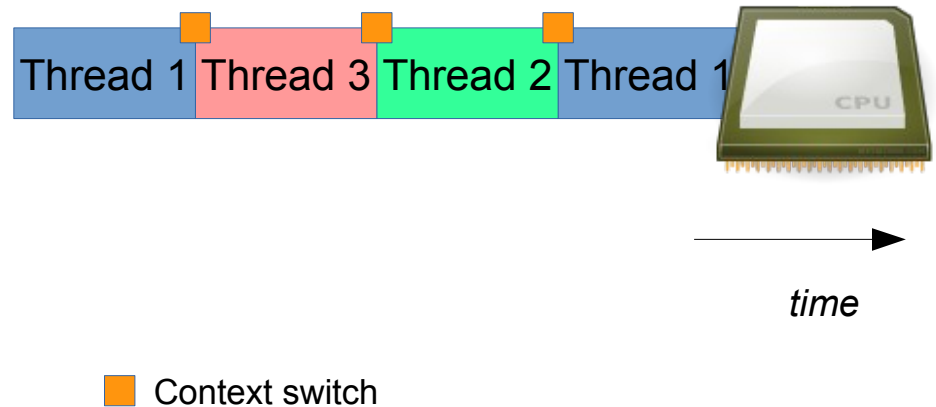
# Execution

- # threads  $\leq$  # cpu cores



Parallel execution

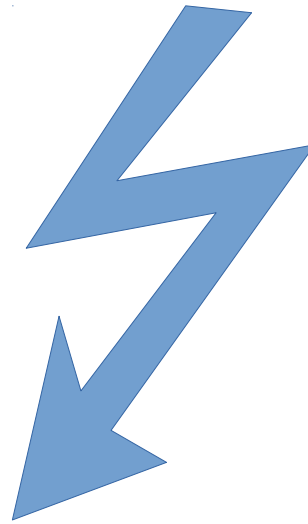
- # threads  $>$  # cpu cores



Pseudo-parallel  
execution

# Implementation

```
void thread1() {  
    // do something  
}  
  
void thread2() {  
    // do something  
}  
  
void main() {  
    while(1) {  
        thread1();  
        thread2();  
    }  
}
```



# Implementation

```
void thread1() {  
    for(int i=0; i<10000; i++) {  
        // Do some computation  
        yield();  
    }  
}  
  
void thread2() {  
    for(int i=0; i<10000; i++) {  
        // Do some computation  
        yield();  
    }  
}  
  
void main() {  
    while(1) {  
        thread1();  
        thread2();  
    }  
}
```

# Yield

```
void thread1() {  
    for(int i=0; i<10000; i++) {  
        // Do some computation  
        yield();  
    }  
}
```



Save context



```
void thread2() {  
    for(int i=0; i<10000; i++) {  
        // Do some computation  
        yield();  
    }  
}
```

Restore context

```
void main() {  
    while(1) {  
        thread1();  
        thread2();  
    }  
}
```



# Yield

```
void thread1() {  
    for(int i=0; i<10000; i++) {  
        // Do some computation  
        yield();  
    }  
}
```

```
void thread2() {  
    for(int i=0; i<10000; i++) {  
        // Do some computation  
        yield();  
    }  
}
```

```
void main() {  
    while(1) {  
        thread1();  
        thread2();  
    }  
}
```

Restore context

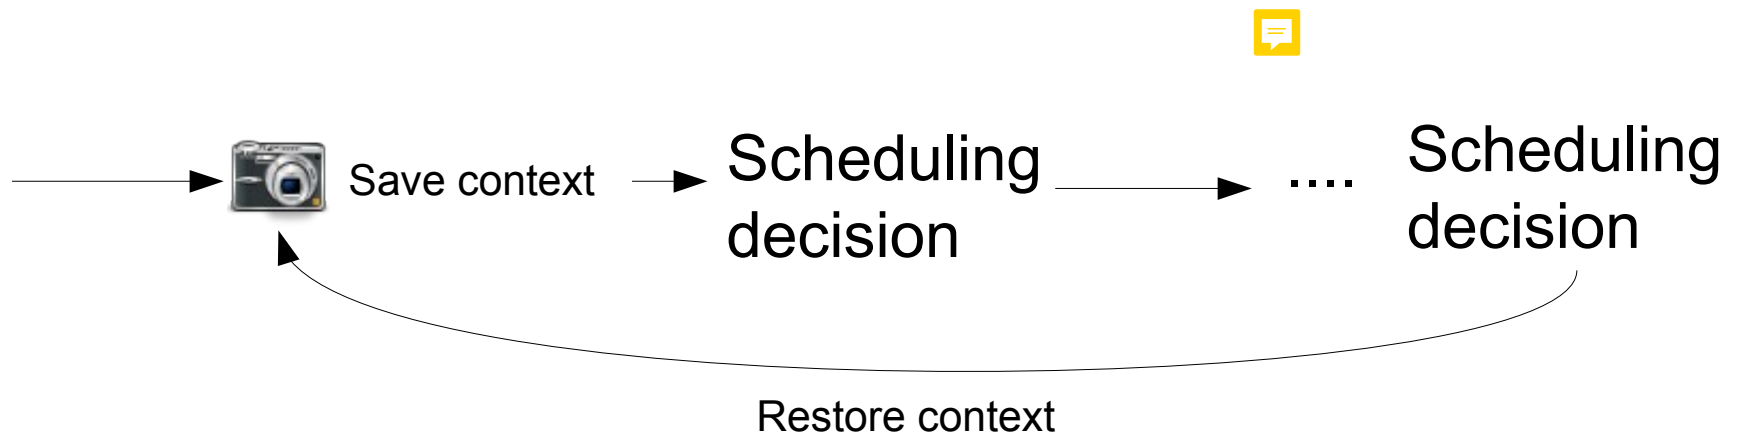


Save context




## Scheduling threads

- To move to another process/thread we must first **save the execution context** of the current process/thread, in order to be able to **restore it later**
  - It's like taking a snapshot of the execution context



## What's the context?

- Register values 
  - General registers
  - Base pointer (“**Where are local variables?**”)
  - Stack pointer (“**Where is the top of the stack?**”)
  - Program counter (“**Where are we in the code?**”)

## Longjmp, setjmp

```
#include <setjmp.h>
```



```
// Save context
```

```
int setjmp(jmp_buf env);
```

```
int sigsetjmp(sigjmp_buf env, int savesigs);
```

```
// Restore context
```

```
void longjmp(jmp_buf env, int val);
```

```
void siglongjmp(sigjmp_buf env, int val);
```

## Example 3.2.1

```
#include <stdio.h>
#include <setjmp.h>

static jmp_buf buf;

void dothings() {
    printf("Now I'm here\n");
    sleep(3);
    longjmp(buf, 42);
    printf("This is never printed\n");
}

int main() {
    if (!setjmp(buf)) { // the first time returns 0
        dothings();
    } else {
        printf("Now I'm there\n");
    }
    return 0;
}
```

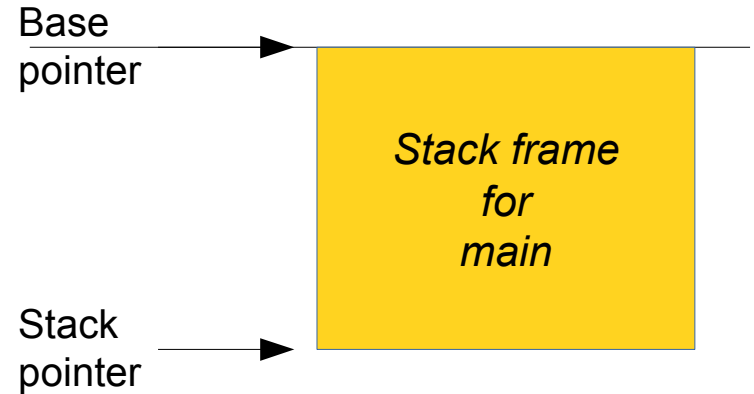
## Example 3.2.1

```
#include <stdio.h>
#include <setjmp.h>
```

```
static jmp_buf buf;
```

```
void dothings() {
    printf("Now I'm here\n");
    sleep(3);
    longjmp(buf, 42);
    printf("This is never printed\n");
}
```

```
int main() {
    if (!setjmp(buf)) {
        dothings();
    } else {
        printf("Now I'm there\n");
    }
    return 0;
}
```



Saves the context into 'buf', returns 0

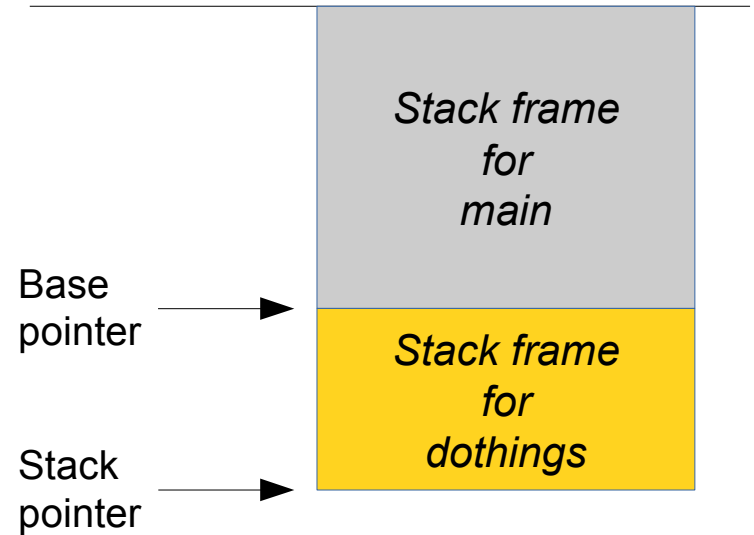
## Example 3.2.1

```
#include <stdio.h>
#include <setjmp.h>
```

```
static jmp_buf buf;
```

```
void dothings() {
    printf("Now I'm here\n");
    sleep(3);
    longjmp(buf, 42);
    printf("This is never printed\n");
}
```

```
int main() {
    if (!setjmp(buf)) { // the first time returns 0
        dothings();
    } else {
        printf("Now I'm there\n");
    }
    return 0;
}
```



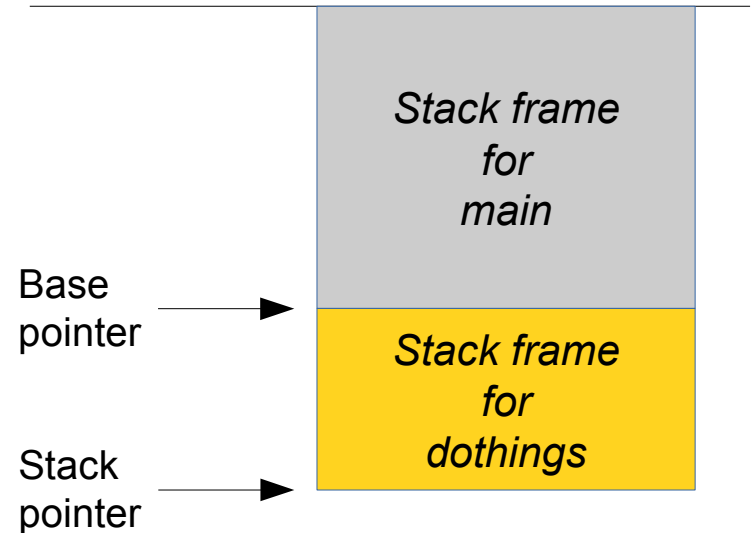
## Example 3.2.1

```
#include <stdio.h>
#include <setjmp.h>
```

```
static jmp_buf buf;
```

```
void dothings() {
    printf("Now I'm here\n");
    sleep(3);
    longjmp(buf, 42);
    printf("This is never printed\n");
}
```

```
int main() {
    if (!setjmp(buf)) { // the first time returns 0
        dothings();
    } else {
        printf("Now I'm there\n");
    }
    return 0;
}
```





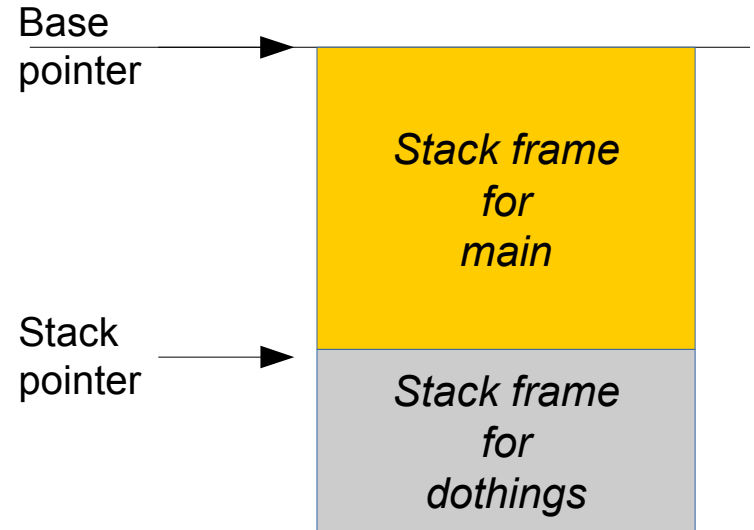
## Example 3.2.1

```
#include <stdio.h>
#include <setjmp.h>
```

```
static jmp_buf buf;
```

```
void dothings() {
    printf("Now I'm here\n");
    sleep(3);
    longjmp(buf, 42);
    printf("This is never printed\n");
}
```

```
int main() {
    if (!setjmp(buf)) { // the first time returns 0
        dothings();
    } else {
        printf("Now I'm there\n");
    }
    return 0;
}
```



Restores the context from 'buf',  
tells 'setjmp' to return 42

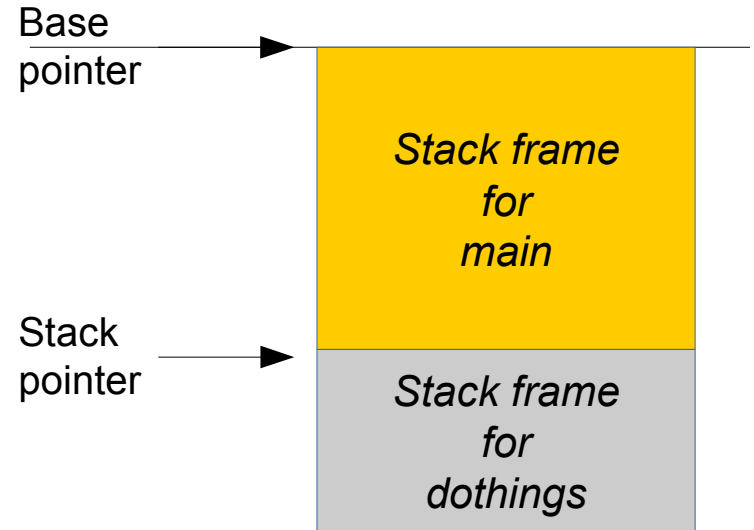
## Example 3.2.1

```
#include <stdio.h>
#include <setjmp.h>
```

```
static jmp_buf buf;
```

```
void dothings() {
    printf("Now I'm here\n");
    sleep(3);
    longjmp(buf, 42);
    printf("This is never printed\n");
}
```

```
int main() {
    if (!setjmp(buf)) { // the first time returns 0
        dothings();
    } else {
        printf("Now I'm there\n");
    }
    return 0;
}
```



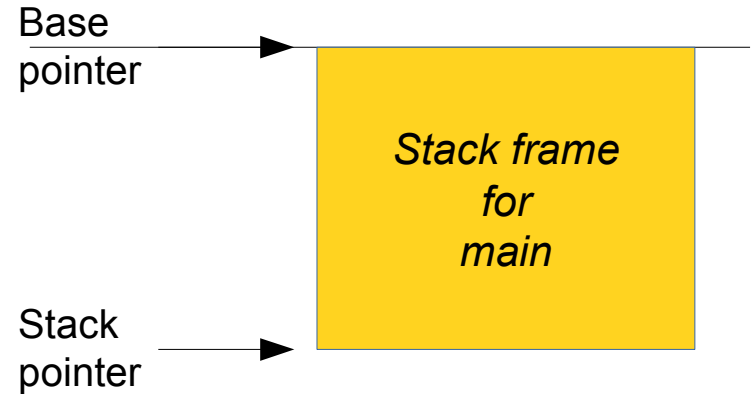
## Example 3.2.2

```
#include <stdio.h>
#include <setjmp.h>
#include <stdlib.h>
```

```
static jmp_buf main_buf, dothings_buf;
```

```
void dothings() {
    int z = 1313;
    if(!setjmp(dothings_buf)) {
        printf("Now I'm here, z=%d\n", z);
        longjmp(main_buf, 42);
    } else {
        printf("Now I'm back here, z=%d\n", z);
        exit(0);
    }
}
```

```
int main() {
    if (!setjmp(main_buf)) {
        dothings();
    } else {
        longjmp(dothings_buf, 17);
    }
    return 0;
}
```



## Example 3.2.2

```
#include <stdio.h>
#include <setjmp.h>
#include <stdlib.h>
```

```
static jmp_buf main_buf, dothings_buf;
```

```
void dothings() {
```

```
    int z = 1313;
```

```
    if(!setjmp(dothings_buf)) {
```

```
        printf("Now I'm here, z=%d\n", z);
```

```
        longjmp(main_buf, 42);
```

```
    } else {
```

```
        printf("Now I'm back here, z=%d\n", z);
```

```
        exit(0);
```

```
    }
```

```
}
```

```
int main() {
```

```
    if (!setjmp(main_buf)) {
```

```
        dothings();
```

```
    } else {
```

```
        longjmp(dothings_buf, 17);
```

```
    }
```

```
    return 0;
```

```
}
```

Base  
pointer

Stack  
pointer

Stack frame  
for  
main

Stack frame Z  
for  
dothings



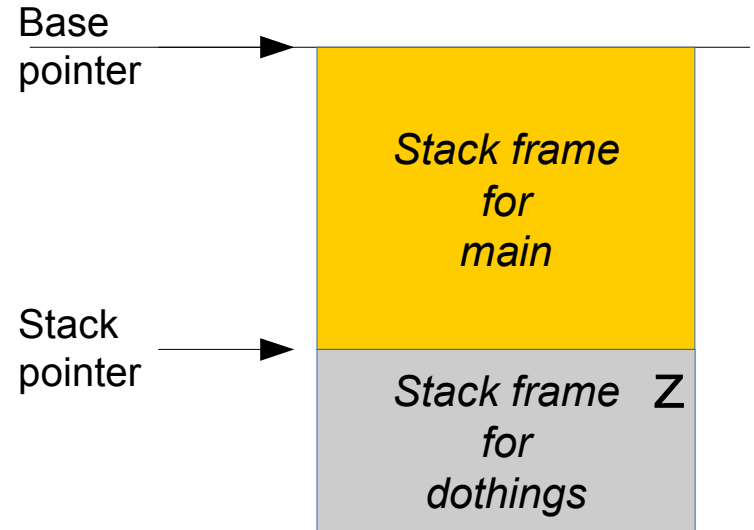
## Example 3.2.2

```
#include <stdio.h>
#include <setjmp.h>
#include <stdlib.h>
```

```
static jmp_buf main_buf, dothings_buf;
```

```
void dothings() {
    int z = 1313;
    if(!setjmp(dothings_buf)) {
        printf("Now I'm here, z=%d\n", z);
        longjmp(main_buf, 42);
    } else {
        printf("Now I'm back here, z=%d\n", z);
        exit(0);
    }
}
```

```
int main() {
    if (!setjmp(main_buf)) {
        dothings();
    } else {
        longjmp(dothings_buf, 17);
    }
    return 0;
}
```



## Example 3.2.2 (stack smashing)

```
#include <stdio.h>
#include <setjmp.h>
#include <stdlib.h>
```

```
static jmp_buf main_buf, dothings_buf;
```

```
void dothings() {
    int z = 1313;
    if(!setjmp(dothings_buf)) {
        printf("Now I'm here, z=%d\n", z);
        longjmp(main_buf, 42);
    } else {
        printf("Now I'm back here, z=%d\n", z);
        exit(0);
    }
}
```

```
int main() {
    if (!setjmp(main_buf)) {
        dothings();
    } else {
        longjmp(dothings_buf, 17);
    }
    return 0;
}
```

Base  
pointer

Stack  
pointer

Stack frame  
for  
main

Stack frame  
for  
dothings

longjmp is also a  
function!

## Example 3.2.2

```
#include <stdio.h>
#include <setjmp.h>
#include <stdlib.h>
```

```
static jmp_buf main_buf, dothings_buf;
```

```
void dothings() {
    int z = 1313;
    if(!setjmp(dothings_buf)) {
        printf("Now I'm here, z=%d\n", z);
        longjmp(main_buf, 42);
    } else {
        printf("Now I'm back here, z=%d\n", z);
        exit(0);
    }
}
```

```
int main() {
    if (!setjmp(main_buf)) {
        dothings();
    } else {
        longjmp(dothings_buf, 17);
    }
    return 0;
}
```

Base  
pointer

Stack  
pointer

Stack frame  
for  
main

????????????????  
????????????????

dothings



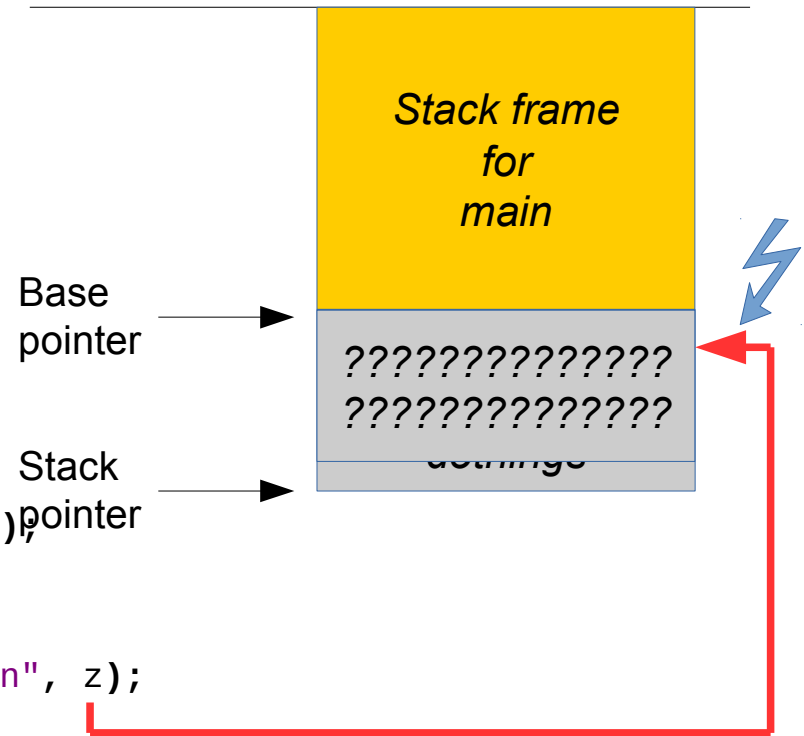
## Example 3.2.2

```
#include <stdio.h>
#include <setjmp.h>
#include <stdlib.h>
```

```
static jmp_buf main_buf, dothings_buf;
```

```
void dothings() {
    int z = 1313;
    if(!setjmp(dothings_buf)) {
        printf("Now I'm here, z=%d\n", z);
        longjmp(main_buf, 42);
    } else {
        printf("Now I'm back here, z=%d\n", z);
        exit(0);
    }
}
```

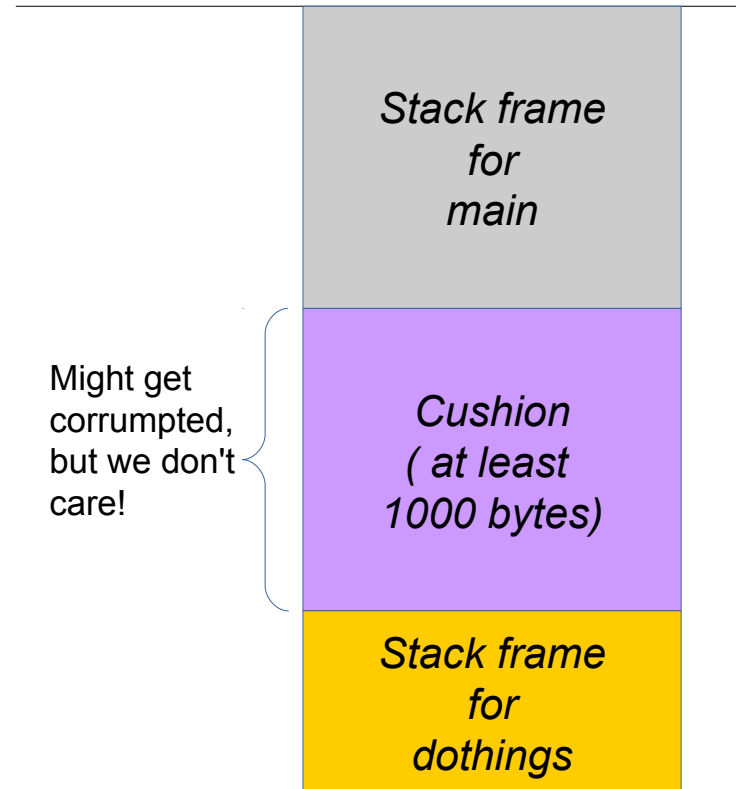
```
int main() {
    if (!setjmp(main_buf)) {
        dothings();
    } else {
        longjmp(dothings_buf, 17);
    }
    return 0;
}
```





# Cushion

```
#include <stdio.h>
#include <setjmp.h>
#include <stdlib.h>
static jmp_buf main_buf, dothings_buf;
void dothings() {
    int z = 1313;
    if(!setjmp(dothings_buf)) {
        printf("Now I'm here, z=%d\n", z);
        longjmp(main_buf, 42);
    } else {
        printf("Now I'm back here, z=%d\n", z);
        exit(0);
    }
}
void cushion() {
    char data_cushion[1000];
    data_cushion[999] = 1;
    dothings();
}
int main() {
    if (!setjmp(main_buf)) {
        cushion();
    } else {
        longjmp(dothings_buf, 17);
    }
    return 0;
}
```



# Cooperative threads example (4)

```
#include <setjmp.h>
#include <stdio.h>

#define CUSHION_SIZE 10000
#define save_context(CONTEXT) setjmp(CONTEXT)
#define restore_context(CONTEXT) longjmp(CONTEXT, 1)

typedef enum { __BTHREAD_UNINITIALIZED, __BTHREAD_READY } bthread_state;
typedef void *(*bthread_routine) (void *);

void create_cushion_and_call(bthread_routine fn, bthread_state* state);
void* bthread1(void* arg);
void* bthread2(void* arg);
void* bthread3(void* arg);

jmp_buf bthread1_buf, bthread2_buf, bthread3_buf;
bthread_state bthread1_state = __BTHREAD_UNINITIALIZED;
bthread_state bthread2_state = __BTHREAD_UNINITIALIZED;
bthread_state bthread3_state = __BTHREAD_UNINITIALIZED;

void create_cushion_and_call(bthread_routine fn, bthread_state* state)
{
    char cushion[CUSHION_SIZE];
    cushion[CUSHION_SIZE-1] = cushion[0];
    *state = __BTHREAD_READY;
    fn(NULL);
}
```



## Cooperative threads example (4)

```
void* bthread1(void* arg)
{
    int i;
    for(i=0;i<10000;i++) {
        printf("BThread1, i=%d\n", i);
        /* Yield to next bthread */
        if (!save_context(bthread1_buf)) {
            if (bthread2_state == __BTHREAD_UNINITIALIZED) {
                create_cushion_and_call(bthread2, &bthread2_state);
            } else {
                restore_context(bthread2_buf);
            }
        }
    }
}

void* bthread2(void* arg)
{
    int i;
    for(i=0;i<10000;i++) {
        printf("BThread2, i=%d\n", i);
        /* Yield to next bthread */
        if (!save_context(bthread2_buf)) {
            if (bthread3_state == __BTHREAD_UNINITIALIZED) {
                create_cushion_and_call(bthread3, &bthread3_state);
            } else {
                restore_context(bthread3_buf);
            }
        }
    }
}
```

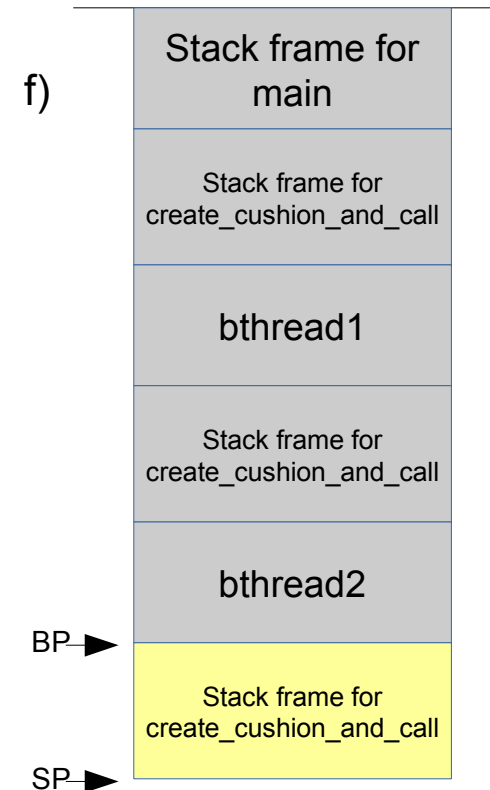
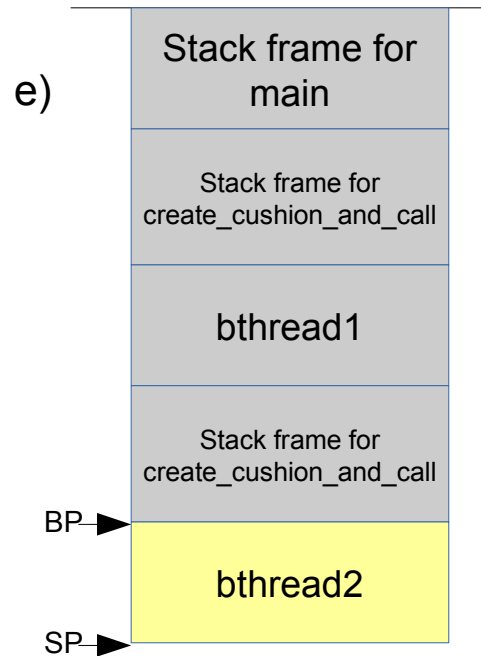
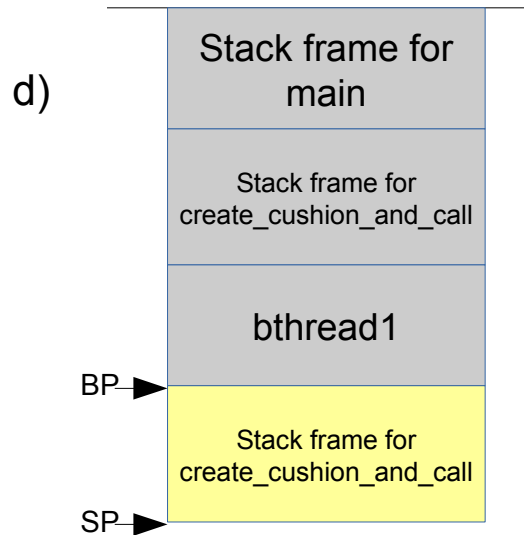
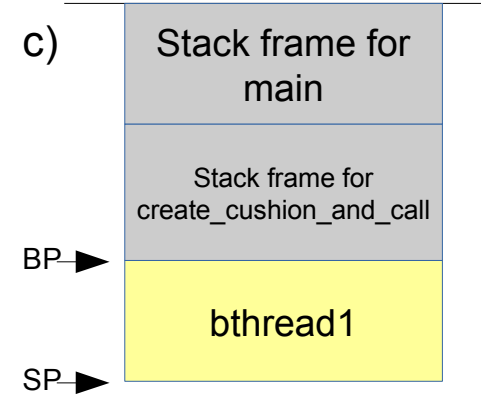
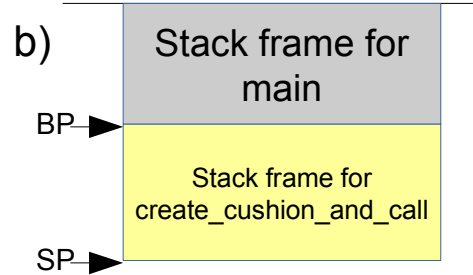
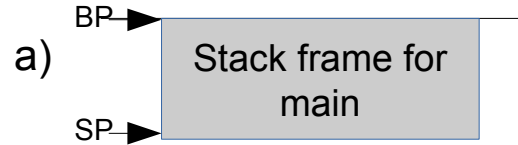
## Cooperative threads example (4)

```
void* bthread3(void* arg)
{
    int i;
    for(i=0;i<10000;i++) {
        printf("BThread3, i=%d\n", i);
        /* Yield to next bthread */
        if (!save_context(bthread3_buf)) {
            // We assume that bthread1 is already
initialized
            restore_context(bthread1_buf, 1);
        }
    }
}

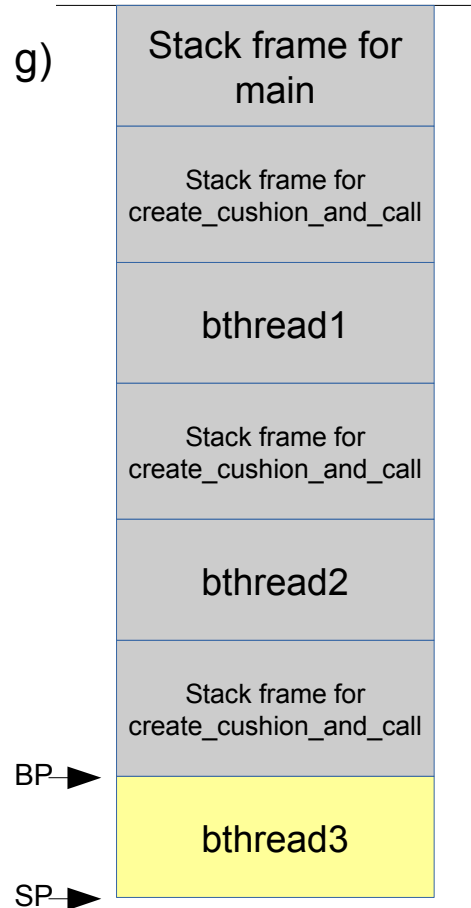
void main()
{
    create_cushion_and_call(bthread1, &bthread1_state);
}
```

BP = Base Pointer, SP = Stack Pointer

## Stack evolution

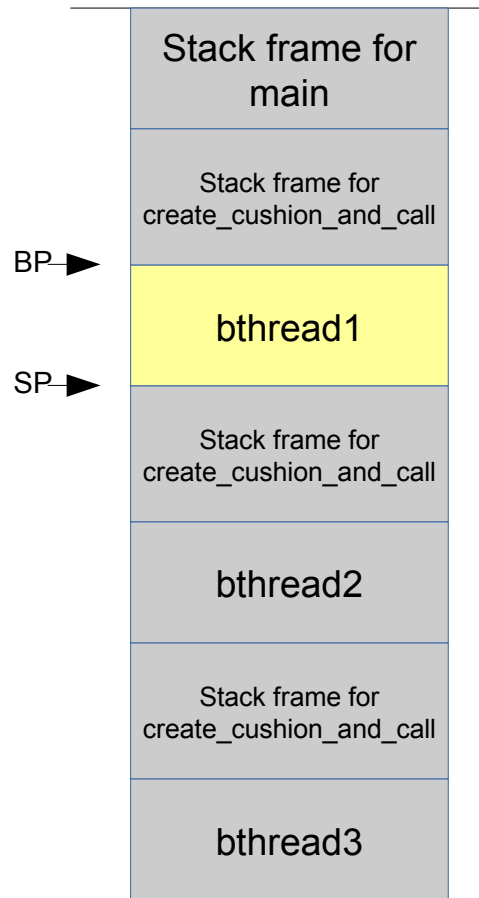


# Stack evolution



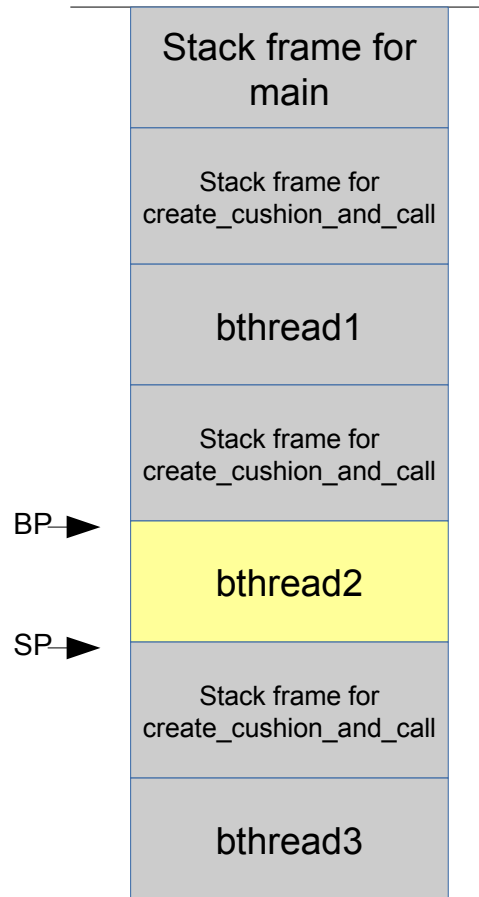
# Context change

Switching (yield) to  
bthread1



# Context change

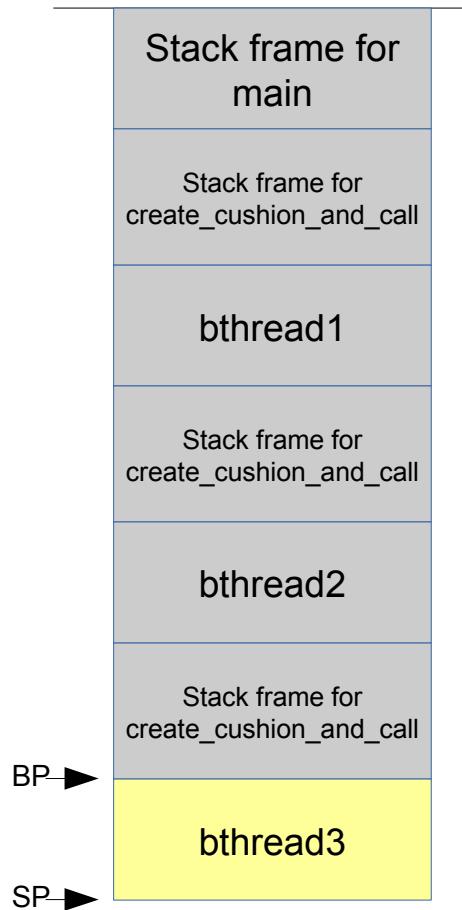
Switching (yield) to  
bthread2





# Context change

Switching (yield) to  
bthread3



## Exercise (2)

### Mandatory

- **Study, compile and test the proposed example code;**
- **Try to add another thread and verify that your code executes as expected.**

Optional (but recommended):

- What happens if a thread exits the for loop earlier than the others?
- How could you solve this problem?
- How could you implement a simple priority mechanism where one receives more CPU time than the other threads?

**... continue reading the tutorial from Section 4.1 !**